

Event driven architecture & Streams

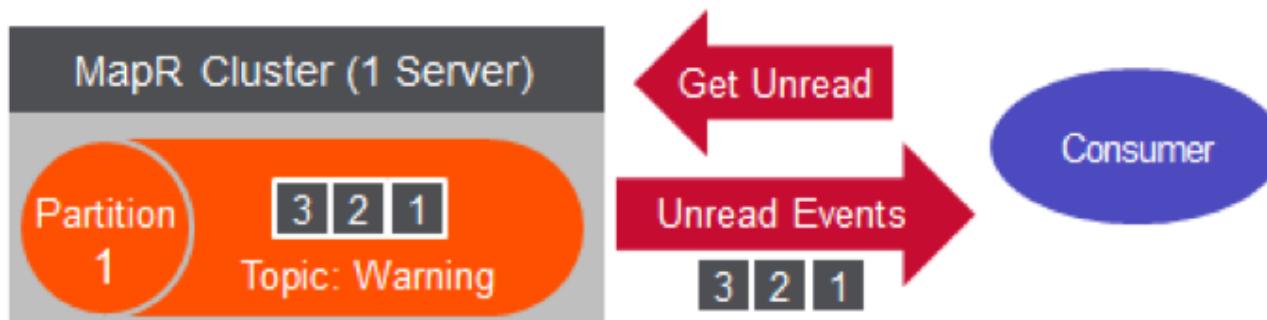
Streams

- Data is not retrieved from a fixed storage device, but rather, handled in near-real-time, with users or other systems rapidly injecting information into our system.
- data injection is asynchronous, where **we do not know ahead of time** when the data will be present.
- In order to facilitate this asynchronous style of data handling, **we have to rethink older polling-based models and instead, use a lighter, more streamlined method.**

Events (message) Stream

- Events are grouped into logical collections of events (often as Topics). Topics can be partitioned for parallel processing. You can think of a partitioned Topic like a queue, **events are delivered in the order they are received.**

Unlike with a queue, events are persisted even after they're delivered



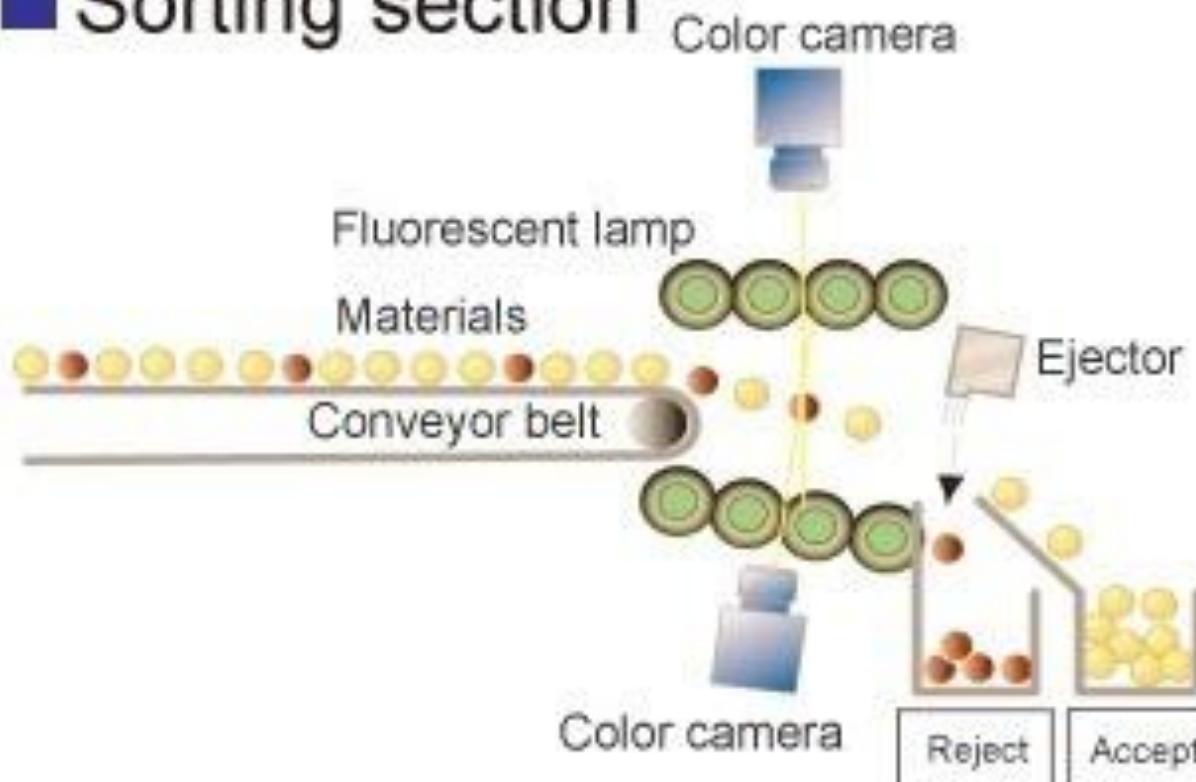
Messages remain on the partition, available to other consumers

Events (message) Stream

- Support multiple consumers
 - topics can have multiple different consumers,
 - remain on the partition, **available to other consumers**.
 - processing of the same messages by different consumers for different purposes.
- Persist events
 - Events are **persisted**, even after they are delivered they
 - Events are **not deleted from Topics when read**
- Can be used to **cache datasets**
 - Older messages are automatically deleted based on the Stream's **time-to-live setting**.
- **Pipelining** is also possible where a consumer enriches an event and publishes it to another topic.

Stream processing (intuition)

■ Sorting section



<https://www.quora.com/What-is-streaming-data>

Event driven architecture models

Pub/sub

- The messaging infrastructure keeps track of subscriptions. When an event is published, it sends the event to each subscriber. After an event is received, it cannot be replayed, and new subscribers do not see the event.

Event streaming

- Events are written to a log. Events are strictly ordered (within a partition) and durable. Clients don't subscribe to the stream, instead a client can read from any part of the stream. The client is responsible for advancing its position in the stream. That means a client can join at any time, and can replay events.

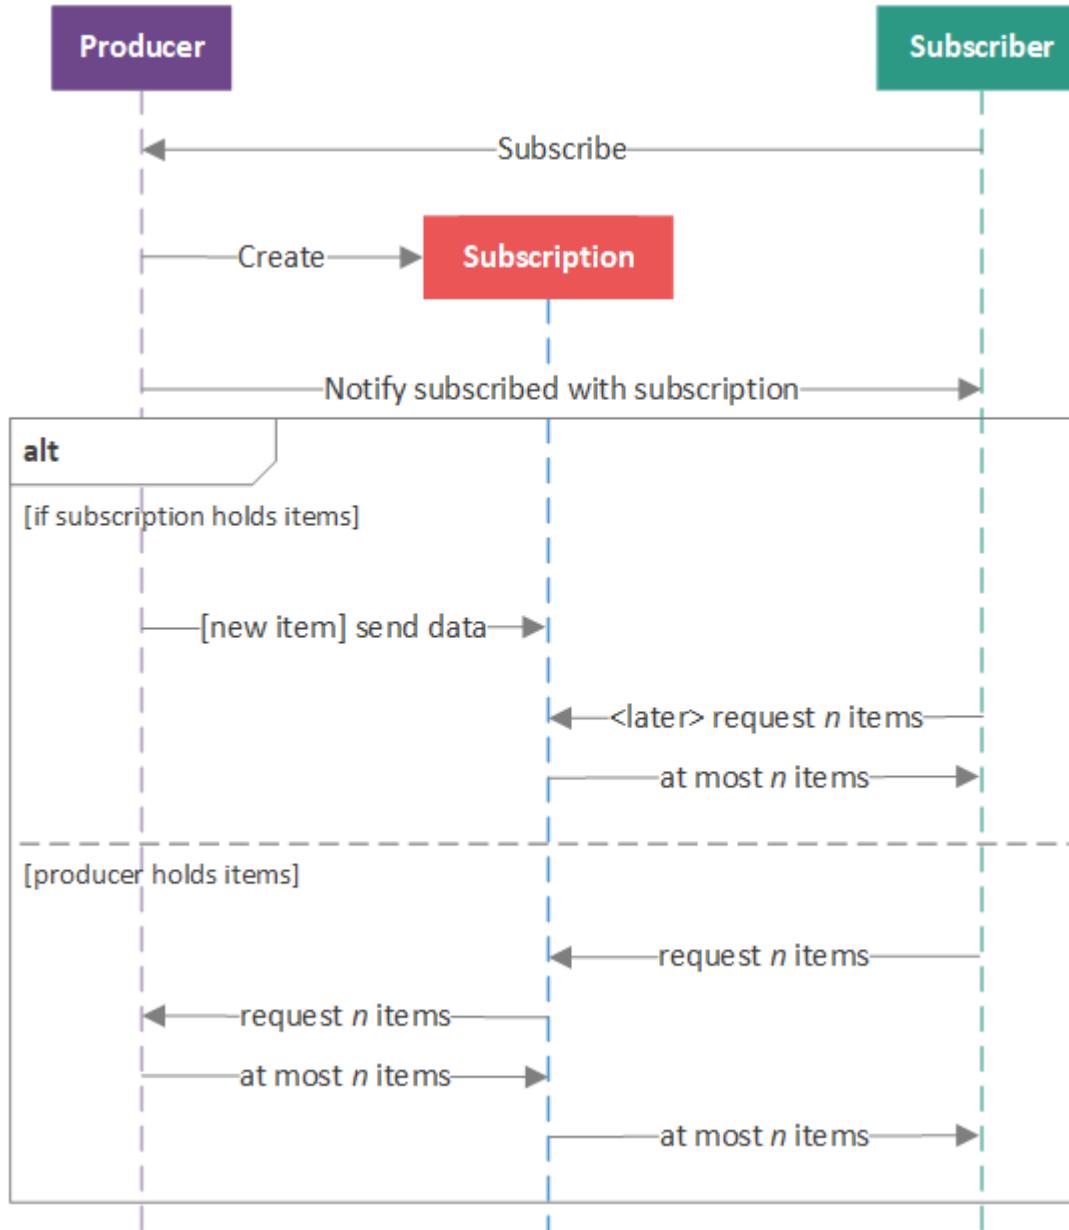
Publishers, Subscribers, and Subscriptions

publisher

- if **items are available**, the publisher pushes the maximum receivable number of items to the subscriber.
- the publisher **pushes** that number of items to the subscriber.

subscriber

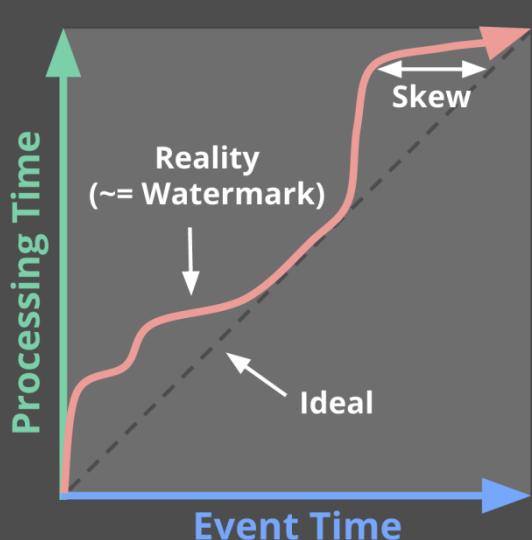
- A **subscriber informs** a **publisher** that it is willing to accept a given number of **items** (**requests** a given number of items)
- the subscriber informs the publisher how many items it is willing to handle



Reactive Streams in Java 9

<https://dzone.com/articles/reactive-streams-in-java-9>

Jumping over “foundations”, hands on



What?	Sums of integers, keyed by team.
Where?	Within fixed event-time windows of one hour.
When?	<ul style="list-style-type: none">• Early: Every 5 minutes of processing time.• On-time: When the watermark passes the end of the window.• Late: Every 10 minutes of processing time.• Final: When the watermark passes the end of the window + two hours.
How?	Panes accumulate new values into prior results.

```
gameEvents
    [... input ...]
    .apply("LeaderboardTeamFixedWindows", Window
        <GameActionInfo>into(FixedWindows.of(
            Duration.standardMinutes(Durations.minutes(60))))
        .triggering(AfterWatermark.pastEndOfWindow()
            .withEarlyfirings(AfterProcessingTime.pastFirstElementInPane()
                .plusDelayOf(Durations.minutes(5)))
            .withLatefirings(AfterProcessingTime.pastFirstElementInPane()
                .plusDelayOf(Durations.minutes(10))))
            .withAllowedLateness(Duration.standardMinutes(120)
                .accumulatingFiredPanes())
        .apply("ExtractTeamScore", new ExtractAndSumScore("team"))
    [... output ...]
```

The world beyond batch: Streaming 101

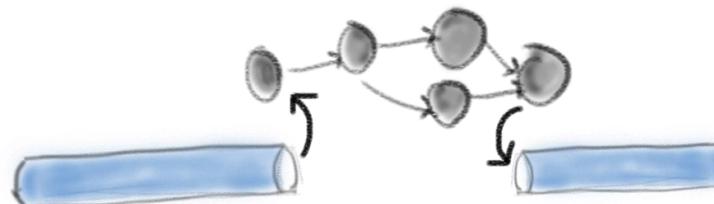
<https://www.oreilly.com/ideas/the-world-beyond-batch-streaming-101>

The world beyond batch: Streaming 102

<https://www.oreilly.com/ideas/the-world-beyond-batch-streaming-102>

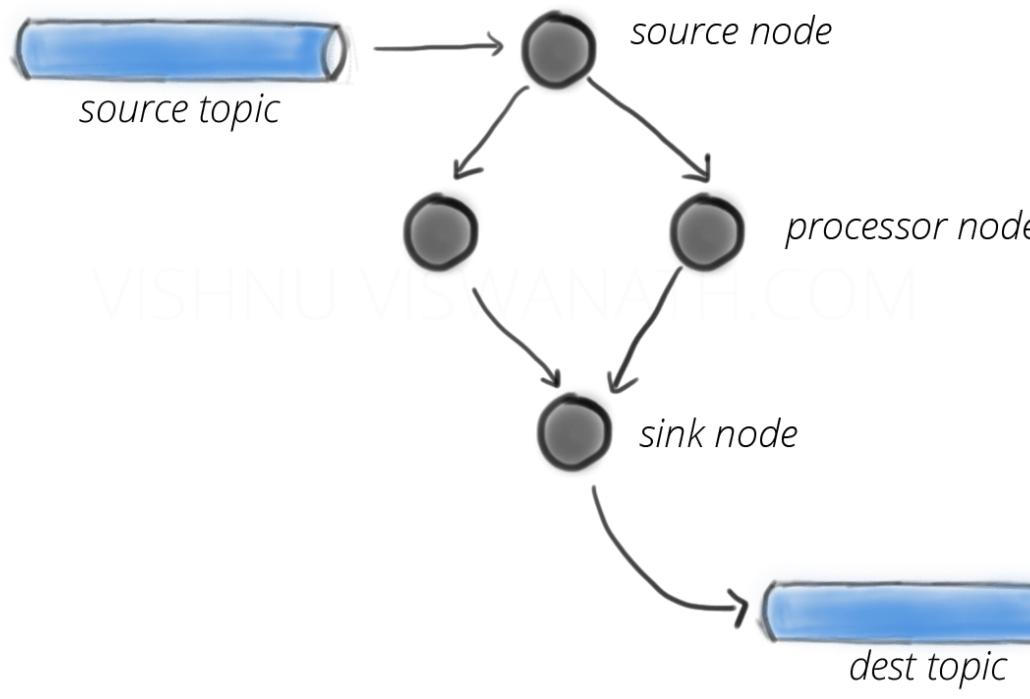


Kafka Streams: an event streaming solution

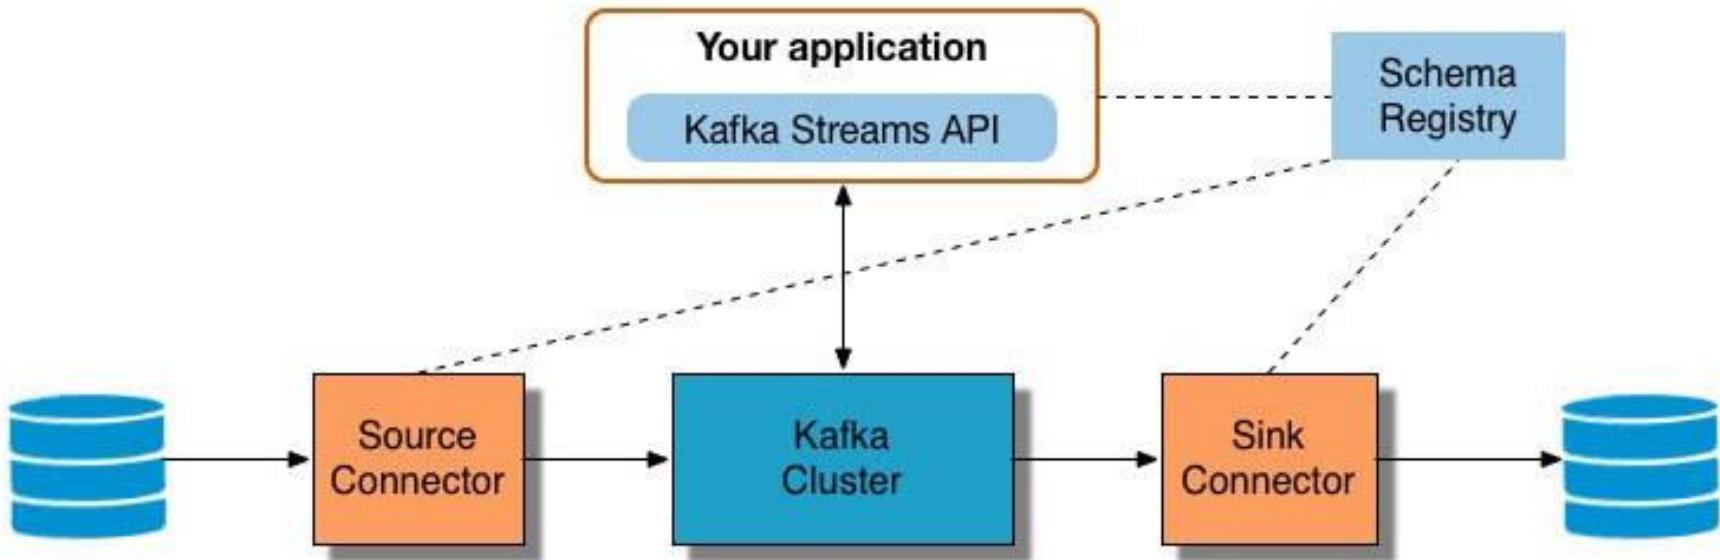


vishnuviswanath.com

Kafka Streams: A newcomer



<http://vishnuviswanath.com/hello-kafka-streams.html>



Building a Real-Time Streaming ETL Pipeline in 20 Minutes
<https://www.confluent.io/blog/building-real-time-streaming-etl-pipeline-20-minutes/>

Obtain a stream from a **stream source**

- collections
- generators
- utility methods
- library classes

Perform one or more **intermediate operations**

- filter
- map
- flatMap
- peek
- distinct
- sorted
- limit
- substream

Black Box

Perform one **terminal operation**

- reducers
- collectors
- forEach
- iterators

<https://stackoverflow.com/questions/22942223/collecting-stream-back-into-the-same-collection-type>

WHITE PAPER

Reactive Programming versus Reactive Systems

Landing on a set of simple Reactive design principles in a sea of constant confusion and overloaded expectations

By Jonas Bonér and Viktor Klang, Lightbend Inc.

25 MIN READ OR DOWNLOAD PDF



Executive Summary

Since co-authoring the [Reactive Manifesto](#) in 2013, we've seen the topic of Reactive go from being a virtually unacknowledged technique for constructing applications—used by only

WHITE PAPER

Reactive Programming versus Reactive Systems

Landing on a set of simple Reactive design principles in a sea of constant confusion and overloaded expectations

When people talk about Reactive in the context of software development and design, they generally mean one of three things:

- Reactive Systems (architecture and design)
- Reactive Programming (declarative event-based)
- Functional Reactive Programming (FRP)

Executive Summary

Since co-authoring the [Reactive Manifesto](#) in 2013, we've seen the topic of Reactive go from being a virtually unacknowledged technique for constructing applications—used by only



Reactive programming

- is not functional or streams
 - Finite observable vs Unbounded stream
 - Can do streams without reactive ...
 - Application scope vs system
- But
 - nice explanation in reactive programming can help
 - Functional paradigm is natural
- “Equivalent” Concepts
 - Pipelines , flow of data
 - Function / processor
 - Source , sink or producer / consumer or observable / consumer

Functional to the rescue

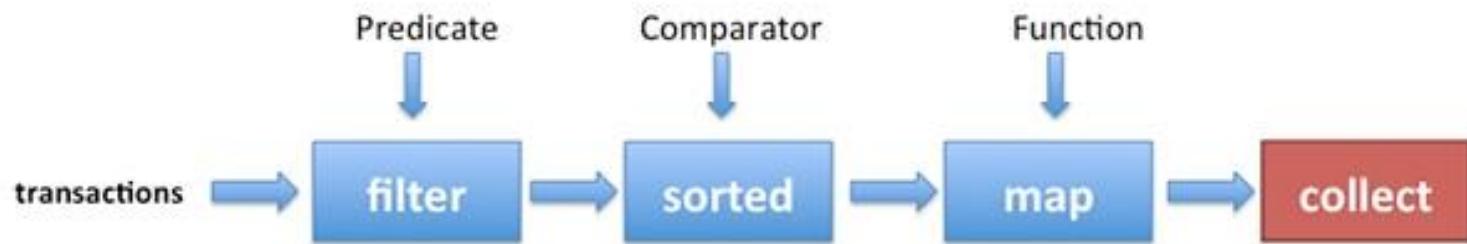
- Loose coupling
 - Functions are “context free”
 - No state
 - Ideally no side effects
 - Can be moved
- Predictable & Scalable
 - load is not dependent on function
 - Runtime, Data
 - Ideally just increase resources
- asynchronous
 - Event driven
 - Stream abstraction
 - Data flow



`map(x => 10 * x)`



java streams are monads

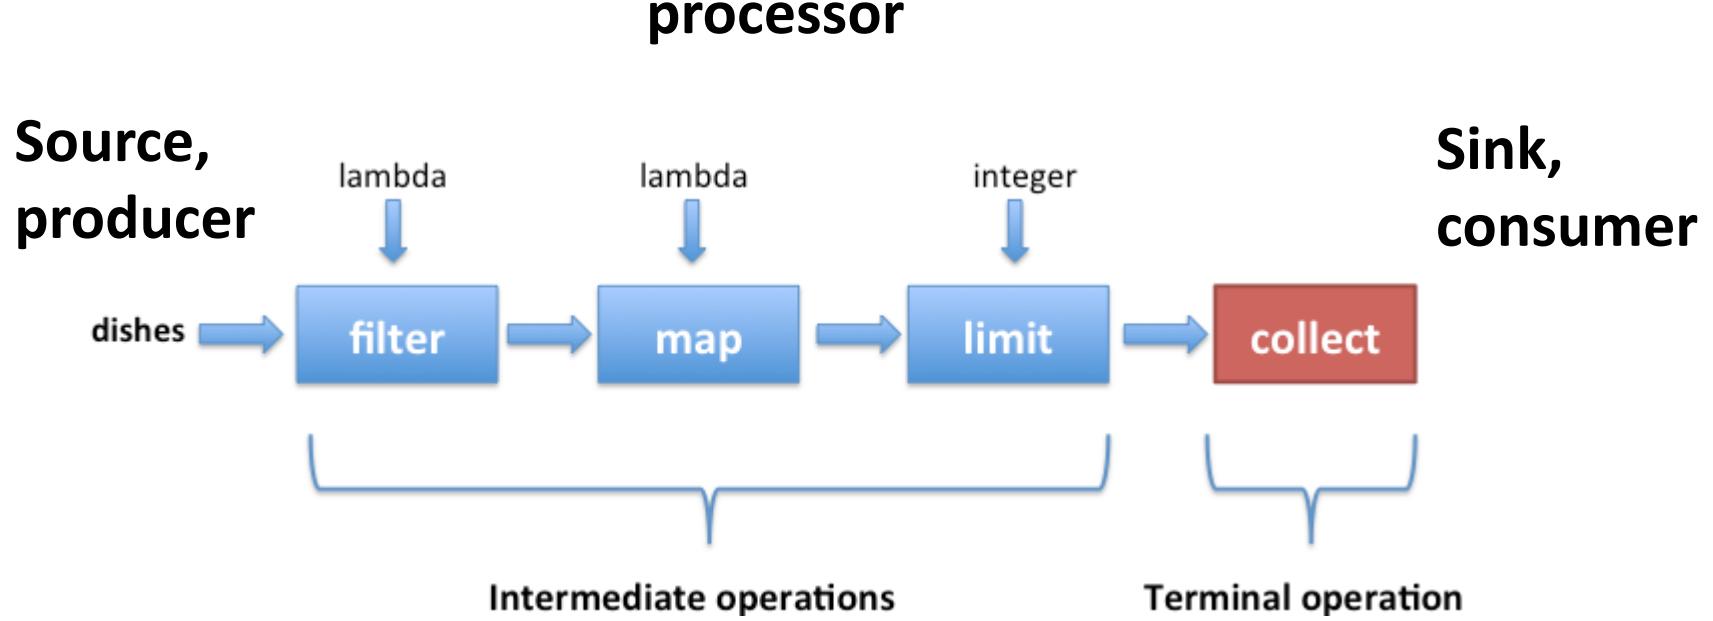


```
List<Integer> transactionIds =  
    transactions.stream()  
        .filter(t -> t.getType() == Transaction.GROCERY)  
        .sorted(comparing(Transaction::getValue).reversed())  
        .map(Transaction::getId)  
        .collect(toList());
```

In functional programming, a monad is a structure that represents computations defined as sequences of steps. A type with a monad structure defines what it means to chain operations, or nest functions of that type together.

Processing Data with Java SE 8 Streams, Part 1

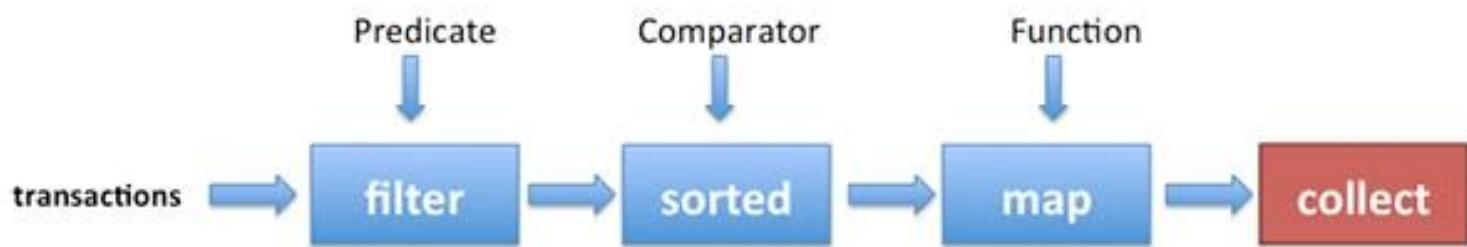
<http://www.oracle.com/technetwork/articles/java/ma14-java-se-8-streams-2177646.html>



Processing Data with Java SE 8 Streams, Part 1

<http://www.oracle.com/technetwork/articles/java/ma14-java-se-8-streams-2177646.html>

java streams are monads



```
List<Integer> transactionsIds =  
    transactions.stream()  
        .filter(t -> t.getType() == Transaction.GROCERY)  
        .sorted(comparing(Transaction::getValue).reversed())  
        .map(Transaction::getId)  
        .collect(toList());
```

In functional programming, a monad is a structure that provides a way to chain computations defined as sequences of functions. The structure defines what it means to chain functions of that type together.

You will see chaining in
“messaging” streams examples
in java

What is a stream?

- is a sequence of records.
- Kafka models a stream as a log, that is, a never-ending sequence of key/value pairs
- often called a changelog

STREAM (as changelog)

("alice",1)

("charlie",1)

("alice", 2)

("bob",1)

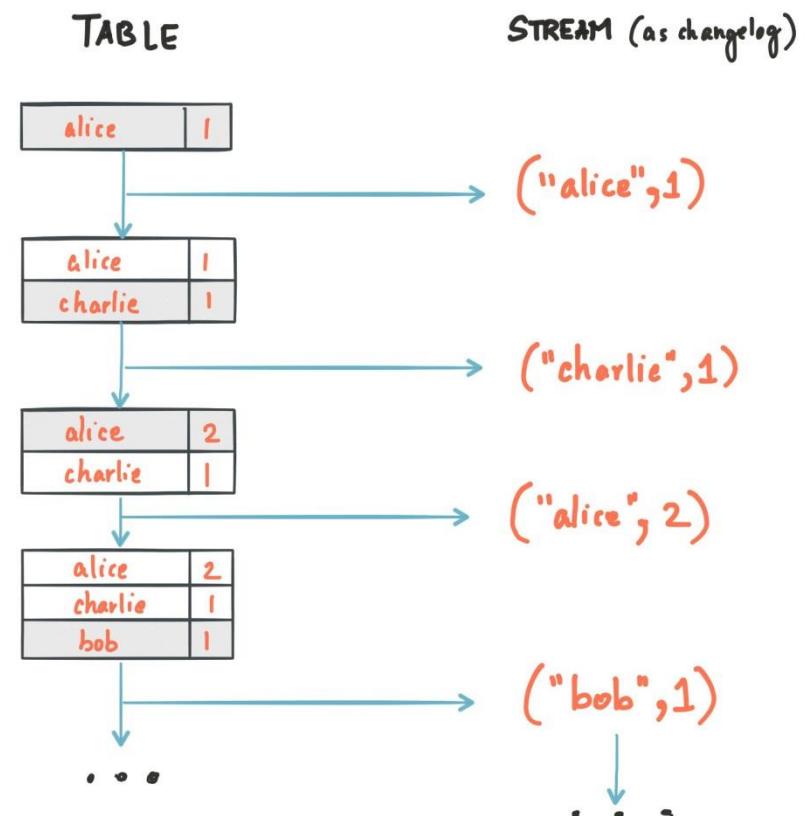
Introducing Kafka Streams: Stream Processing Made Simple

<https://www.confluent.io/blog/introducing-kafka-streams-stream-processing-made-simple/>

<https://docs.confluent.io/currentstreams/concepts.html#kafka-101>

Tables and streams

- a pure stream is one where the changes are interpreted as just INSERT statements (since no record replaces any existing record), whereas a table is a stream where the changes are interpreted as UPDATEs (since any existing row with the same key is overwritten).



Creating source streams from Kafka

To stream

To ktable

```
StreamsBuilder builder = new StreamsBuilder();

KStream<String, Long> wordCounts = builder.stream(
    "word-counts-input-topic", /* input topic */
    Consumed.with(
        Serdes.String(), /* key serde */
        Serdes.Long()   /* value serde */
    );
}
```

<https://docs.confluent.io/currentstreams/developer-guide/dsl-api.html>

How does this relate to streams & tables?

Tables capture a point-in-time snapshot of a time-varying relation.

Streams capture the evolution of a time-varying relation over time.

Foundations of Streaming SQL or: How I learned to love stream & table theory [Strata NYC 2017]

<https://www.youtube.com/watch?v=YO95iAkwapQ>

<https://www.infoq.com/presentations/beam-model-stream-table-theory>

General theory of stream & table relativity

Pipelines : **tables** + **streams** + **operations**

Tables : data at rest

Streams : data in motion

Operations : (**stream** | **table**) → (**stream** | **table**) transformations

- **stream** → **stream**: Non-grouping (element-wise) operations
Leaves stream data in motion, yielding another stream.
- **stream** → **table**: Grouping operations
Brings stream data to rest, yielding a table.
Windowing adds the dimension of time to grouping.
- **table** → **stream**: Ungrouping (triggering) operations
Puts table data into motion, yielding a stream.
Accumulation dictates the nature of the stream (deltas, values, retractions).
- **table** → **table**: (none)
Impossible to go from rest and back to rest without being put into motion.

Foundations of Streaming SQL or: How I learned to love stream & table theory [Strata NYC 2017]

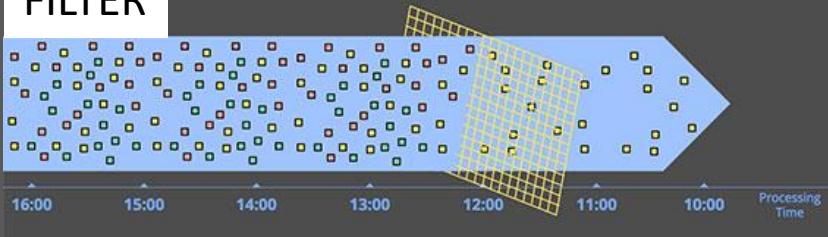
<https://www.youtube.com/watch?v=YO95iAkwapQ>

<https://www.infoq.com/presentations/beam-model-stream-table-theory>

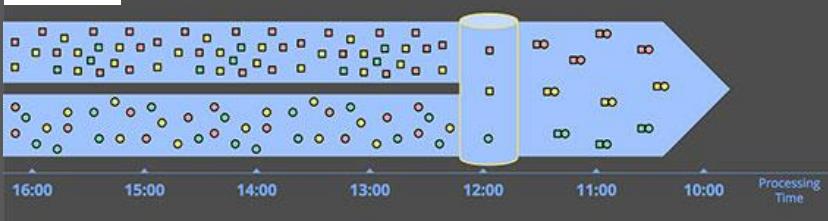


Operations on streams

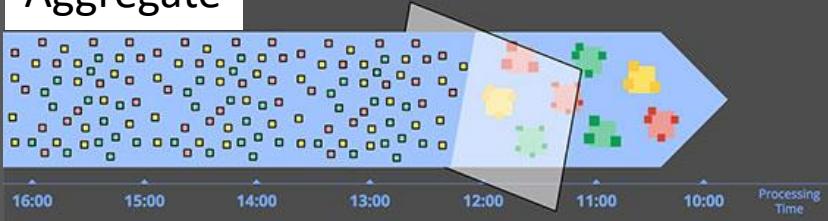
FILTER



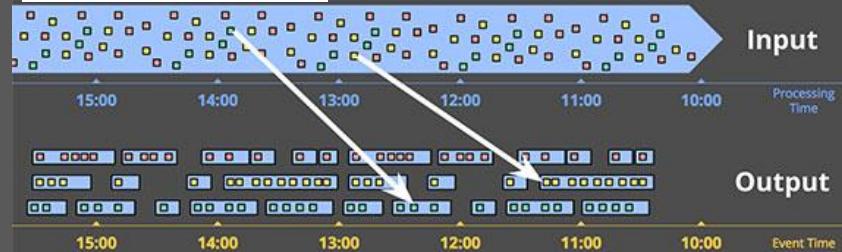
Join



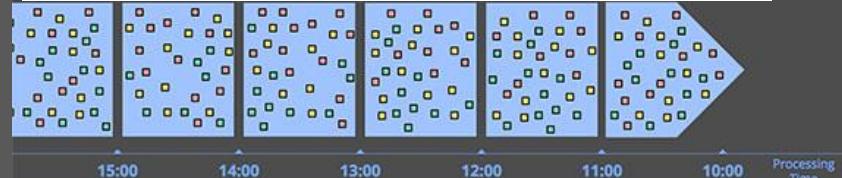
Aggregate



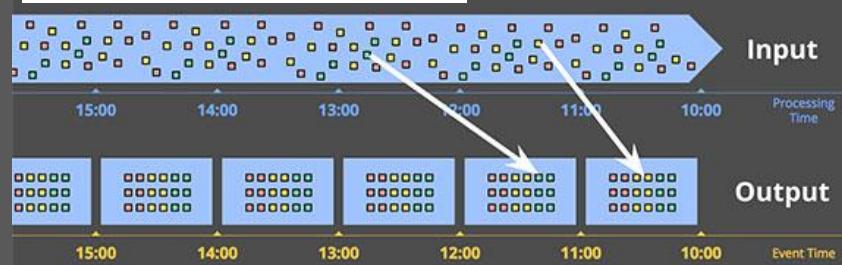
Split by event



Fixed windowing by processing time



Windowing by event



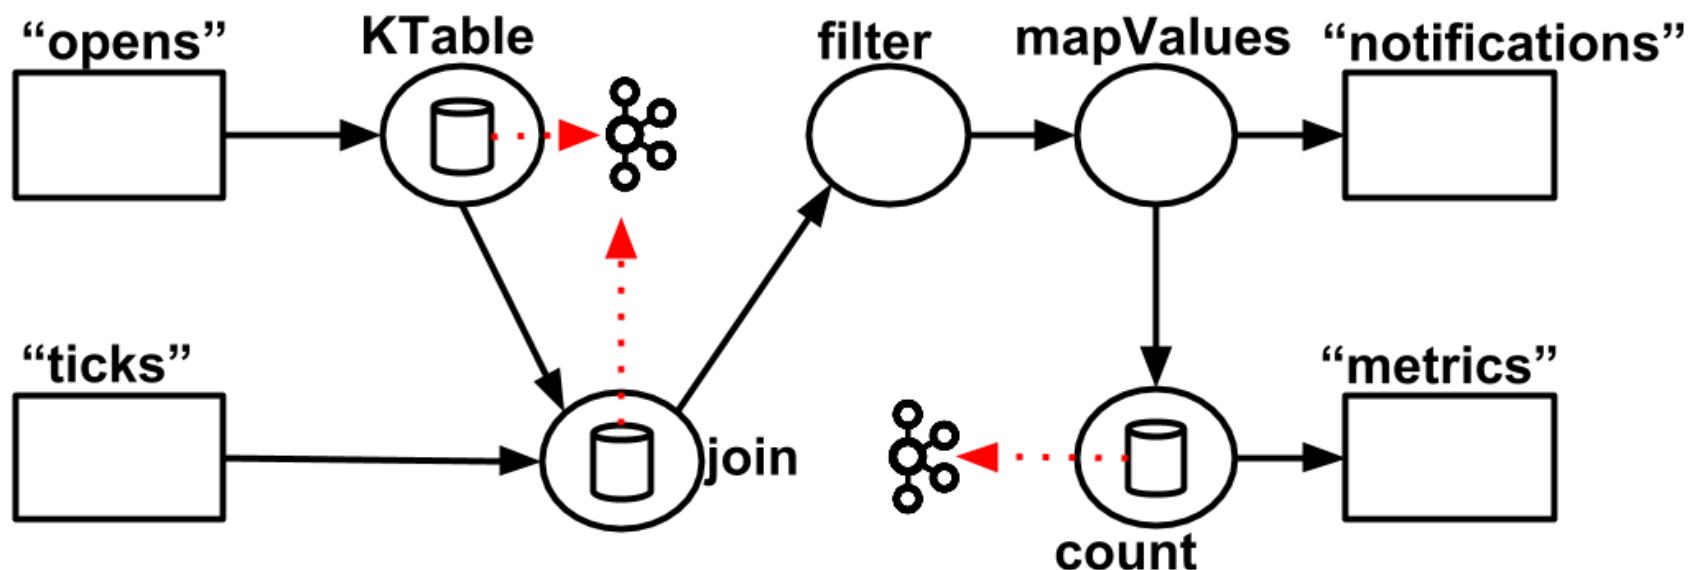
The world beyond batch: Streaming 101

<https://www.oreilly.com/ideas/the-world-beyond-batch-streaming-101>

The world beyond batch: Streaming 102

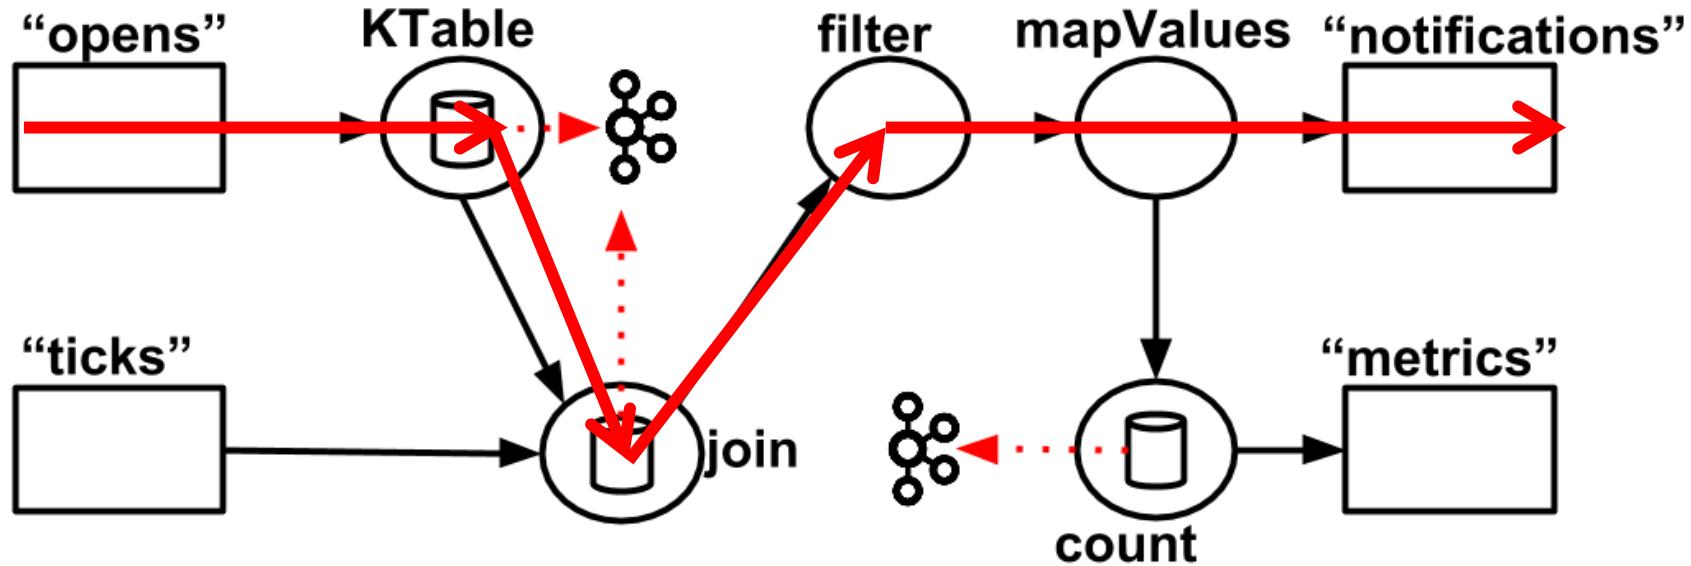
<https://www.oreilly.com/ideas/the-world-beyond-batch-streaming-102>





Kafka Streams - Not Looking at Facebook

<https://timothyrenner.github.io/engineering/2016/08/11/kafka-streams-not-looking-at-facebook.html>



```

// Stream from a Kafka topic.
builder.stream("open")

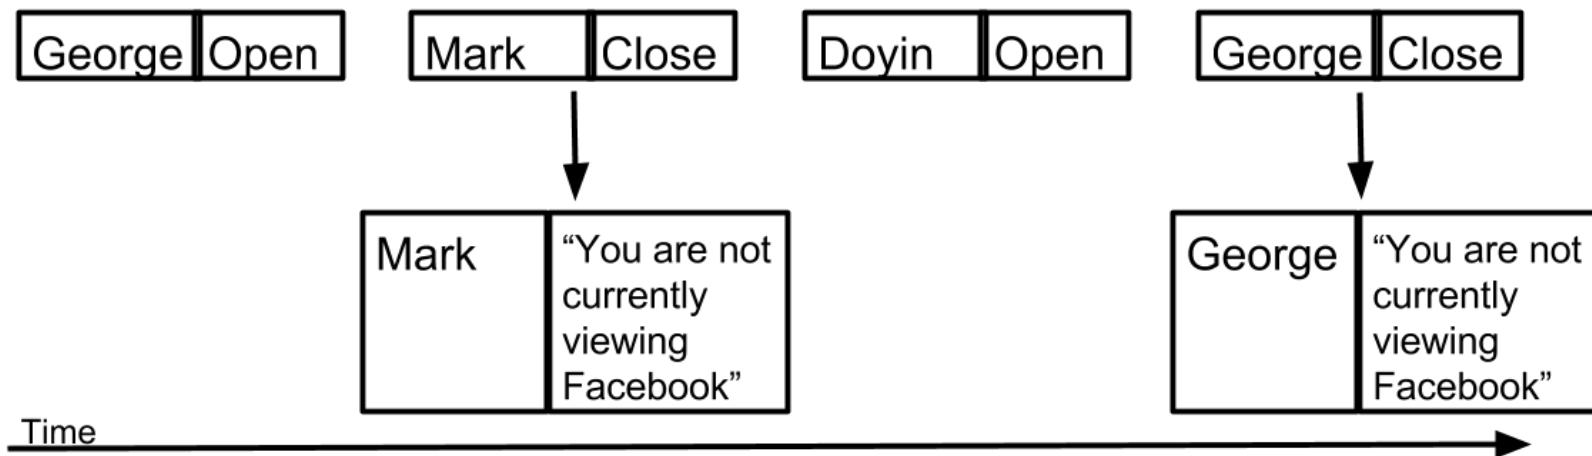
// Remove "Open" events.
.filter((k,v) -> v == "Close")

// Add the helpful notification message.
.mapValues(v -> "You are not currently viewing Facebook.")

// Sink to a Kafka topic.
.to("notifications");

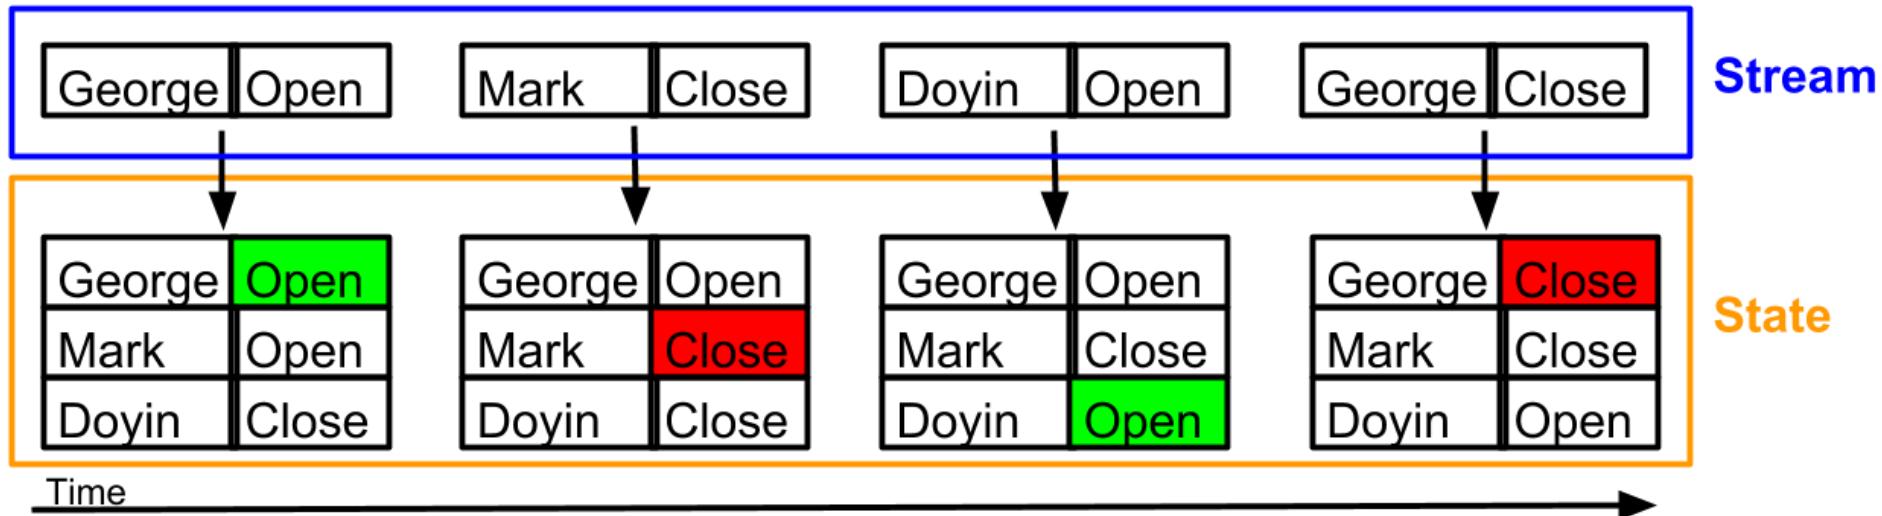
```

KStream

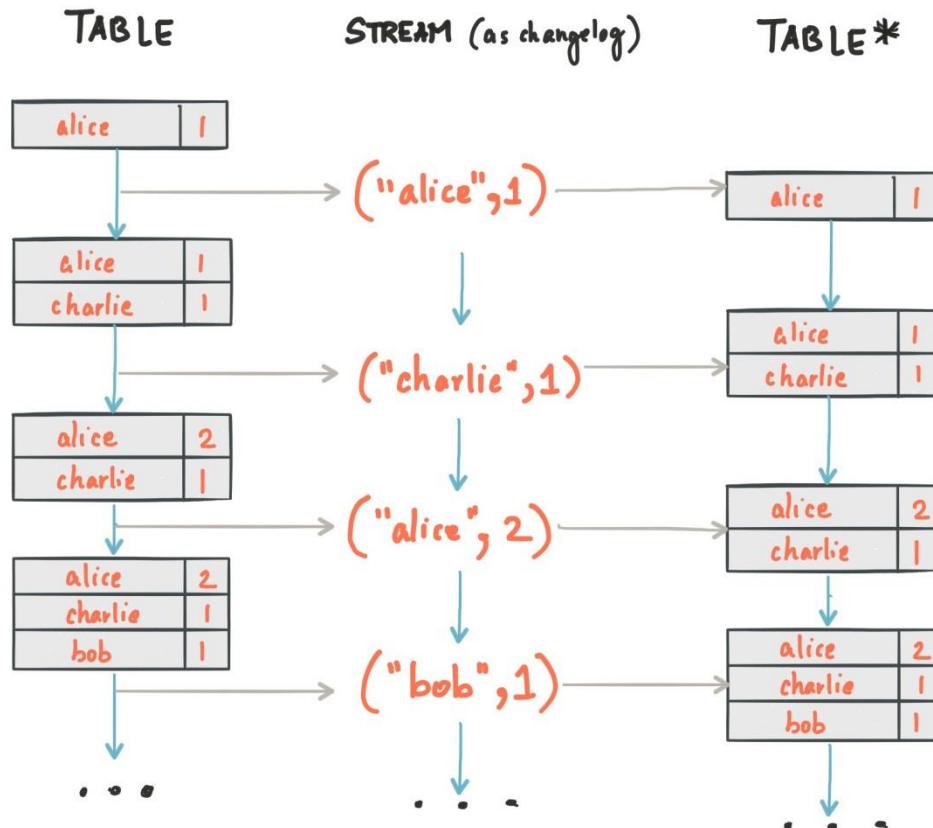


```
// Stream from a Kafka topic.  
builder.stream("open")  
  
// Remove "Open" events.  
.filter((k,v) -> v == "Close")  
  
// Add the helpful notification message.  
.mapValues(v -> "You are not currently viewing Facebook.")  
  
// Sink to a Kafka topic.  
.to("notifications");
```

ktable



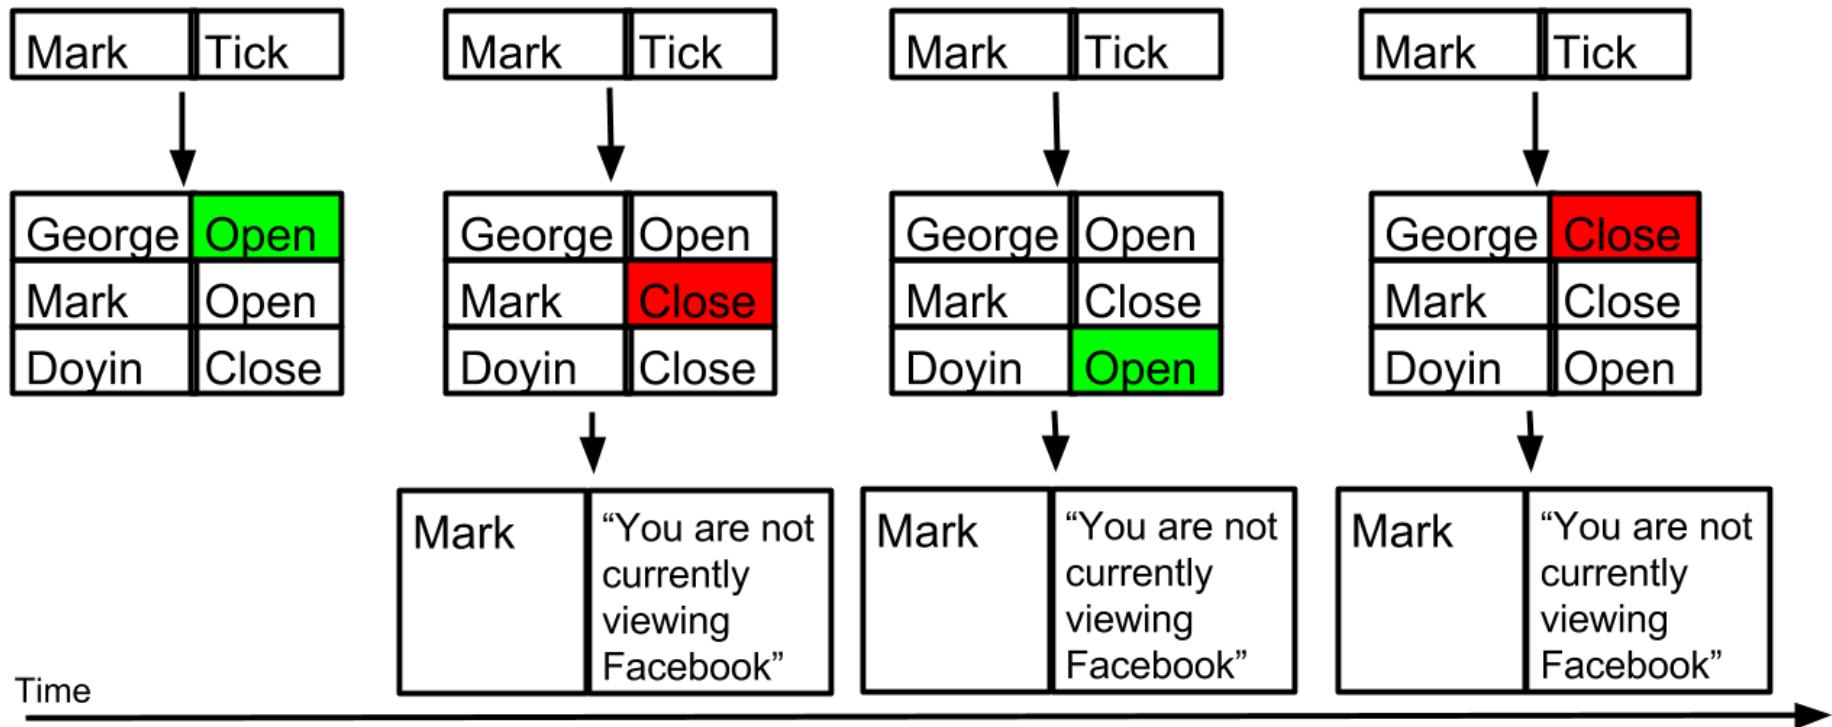
Tables and Streams are Dual

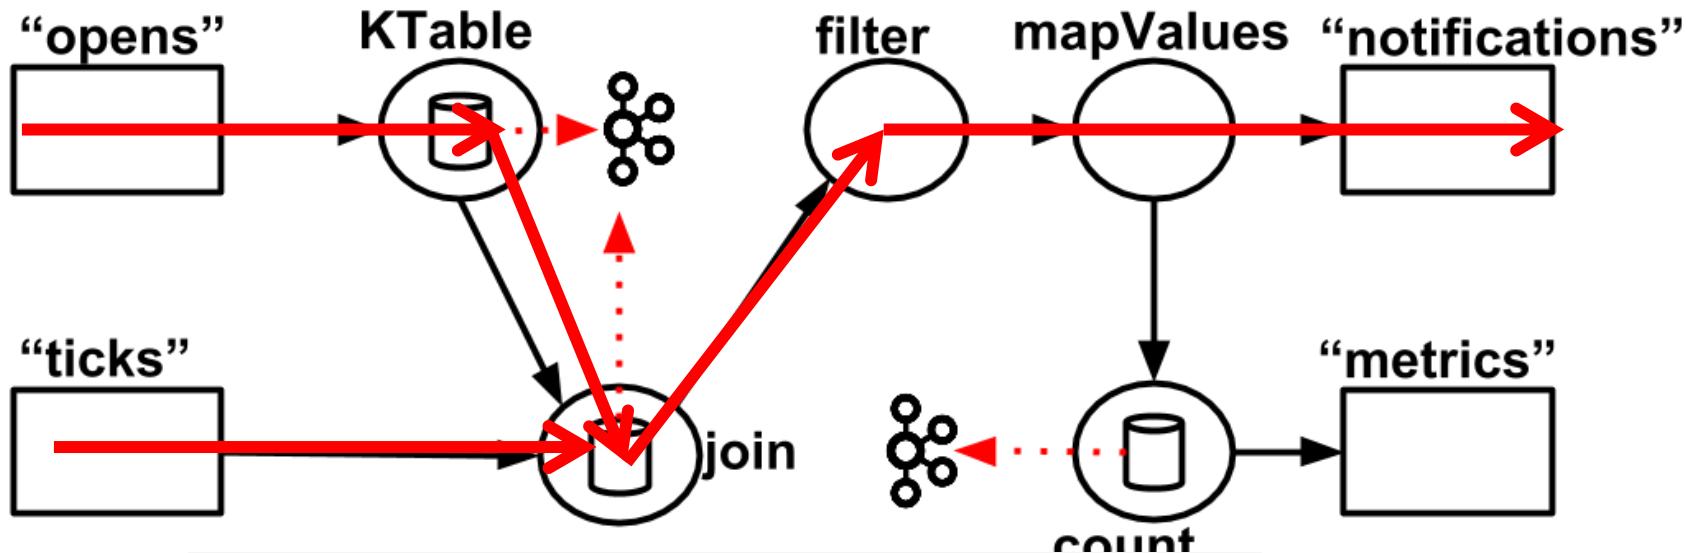


Introducing Kafka Streams: Stream Processing Made Simple

<https://www.confluent.io/blog/introducing-kafka-streams-stream-processing-made-simple/>

Ktable in the middle of streams





```

KStream<String, String> ticks = builder.stream("tick");

// Now we stream "open" as a table.
KTable<String, String> open = builder.table("open");

// Join ticks to the open table.
KStream<String, String> closedTicks =
    ticks.leftJoin(open, (vtick, vopen) -> vopen)

        // Filter the "closed" entries in the new stream.
        .filter((k,v) -> v == "Closed");

        // Now fill in the message.
        .mapValues(v -> "You are not currently viewing Facebook.");

        // Sink the messages to the "notifications" topic.
        closedTicks.to("notifications");
    
```

Ktable in the middle of streams

Mark

George

Mark

Doyin

Time

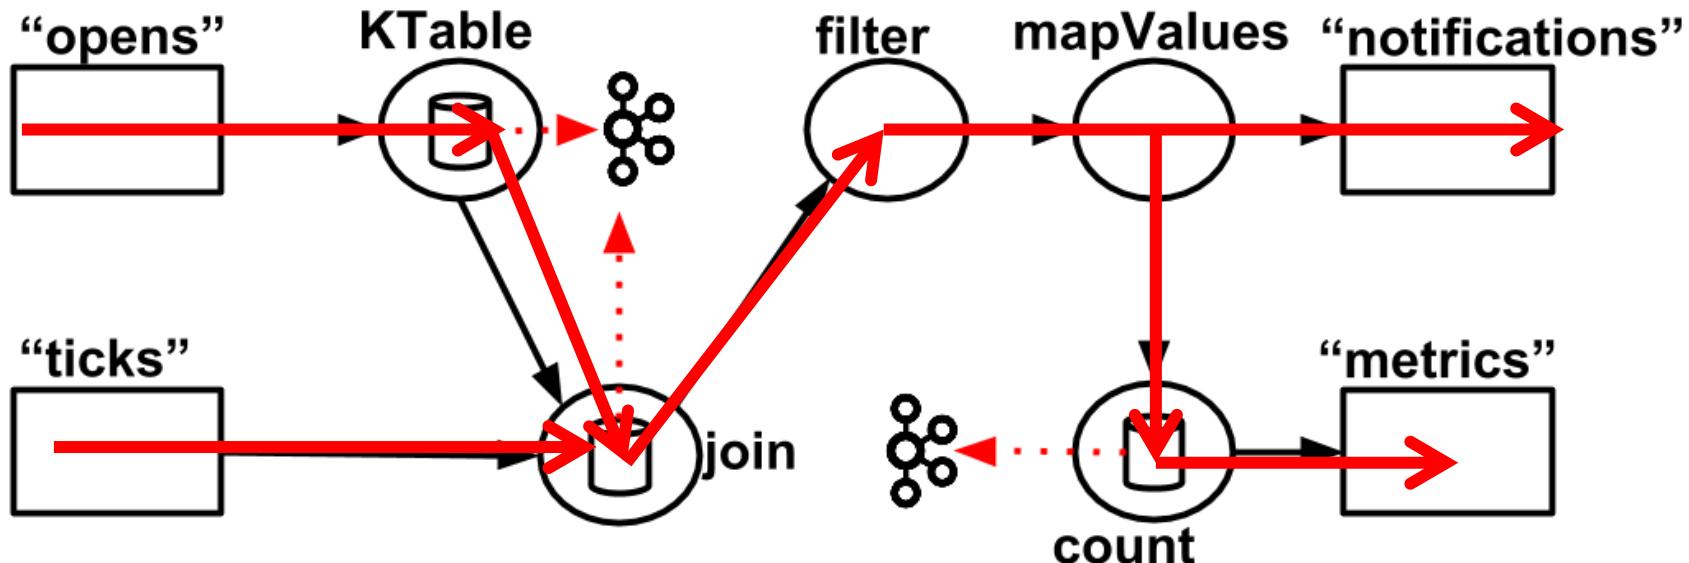
```
KStream<String, String> ticks = builder.stream("tick");

// Now we stream "open" as a table.
KTable<String, String> open = builder.table("open");

// Join ticks to the open table.
KStream<String, String> closedTicks =
    ticks.leftJoin(open, (vtick, vopen) -> vopen)
        // Filter the "closed" entries in the new stream.
        .filter((k,v) -> v == "Closed");

        // Now fill in the message.
        .mapValues(v -> "You are not currently viewing Facebook.");

// Sink the messages to the "notifications" topic.
closedTicks.to("notifications");
```



```

KTable<Windowed<String>, Long> notificationCounts =
    closedTicks.countByKey(
        TimeWindows.of("notificationCounts", 60000L * 10)
        // "Hop" the windows every minute.
        .advanceBy(60000L)
        // Ignore late values.
        .until(60000L * 10));

    // Convert the table updates to a stream.
    // Note we need to extract the key from the window.
    notificationCounts.toStream((k,v) -> k.key())
        .to("metrics");

```

Windows

```
KTable<Windowed<String>, Long> notificationCounts =  
    closedTicks.countByKey(  
        TimeWindows.of("notificationCounts", 60000L * 10)  
            // "Hop" the windows every minute.  
            .advanceBy(60000L)  
            // Ignore late values.  
            .until(60000L * 10));  
  
    // Convert the table updates to a stream.  
    // Note we need to extract the key from the window.  
    notificationCounts.toStream((k,v) -> k.key())  
        .to("metrics");
```

Basic example

```
(...)
public class Stream {
    public static void main(final String[] args) throws Exception {
        final Properties streamsConfiguration = new Properties();
        streamsConfiguration.put(StreamsConfig.APPLICATION_ID_CONFIG, "messageStream");
        streamsConfiguration.put(StreamsConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
        streamsConfiguration.put(StreamsConfig.VALUE_SERDE_CLASS_CONFIG, Serdes.String().getClass().getName());

        final Serde<String> stringSerde = Serdes.String();
        final KStreamBuilder builder = new KStreamBuilder();
        final KStream<String, String> incomingMessageStream = builder.stream(stringSerde, stringSerde, "messages");
        final KStream<String, String> processedStream = incomingMessageStream.mapValues(message -> message + " - Processed");

        processedStream.to(stringSerde, stringSerde, "messagesProcessed");

        final KafkaStreams streams = new KafkaStreams(builder, streamsConfiguration);
        streams.start();
        Runtime.getRuntime().addShutdownHook(new Thread(streams::close));
    }
}
```

How to process streams with Kafka Streams?

<https://medium.com/@yagizdemirsoy/how-to-process-streams-with-kafka-streams-e1dfefe20466>

(des)erialization

```
final Serde<String> stringSerde = Serdes.String();
final Serde<Long> longSerde = Serdes.Long();

KStreamBuilder builder = new KStreamBuilder();
KStream<String, String> textLines = builder.stream(stringSerde, stringSerde, "TextLinesTopic");
KStream<String, Long> wordCounts = textLines
    .flatMapValues(value -> Arrays.asList(value.toLowerCase().split("\\w+")))
    .map((key, word) -> new KeyValue<>(word, word))
    // Required in Kafka 0.10.0 to re-partition the data because we re-keyed the stream in the `map` step.
    // Upcoming Kafka 0.10.1 does this automatically for you (no need for `through`).
    .through("RekeyedIntermediateTopic")
    .countByKey("Counts")
    .toStream();
wordCounts.to(stringSerde, longSerde, "WordsWithCountsTopic");

KafkaStreams streams = new KafkaStreams(builder, streamsConfiguration);
streams.start();

Runtime.getRuntime().addShutdownHook(new Thread(streams::close));
```

- Kafka is content agnostic so it needs:
- deserialization to extract context from message
- serialization to sent to stream

<https://gist.github.com/miguno/a9b44c1efc9b579565cfad75e543c808>

Streams from/to topics

```
final Serde<String> stringSerde = Serdes.String();
final Serde<Long> longSerde = Serdes.Long();

KStreamBuilder builder = new KStreamBuilder();
KStream<String, String> textLines = builder.stream(stringSerde, stringSerde, "TextLinesTopic");
KStream<String, Long> wordCounts = textLines
    .flatMapValues(value -> Arrays.asList(value.toLowerCase().split("\\w+")))
    .map((key, word) -> new KeyValue<>(word, word))
    // Required in Kafka 0.10.0 to re-partition the data because we re-keyed the stream in the `map` step.
    // Upcoming Kafka 0.10.1 does this automatically for you (no need for `through`).
    .through("RekeyedIntermediateTopic")
    .countByKey("Counts")
    .toStream();
wordCounts.to(stringSerde, longSerde, "WordsWithCountsTopic");

KafkaStreams streams = new KafkaStreams(builder, streamsConfiguration);
streams.start();

Runtime.getRuntime().addShutdownHook(new Thread(streams::close));
```

Stream origin are streams and generate streams – in kafka these are placed in topics

<https://gist.github.com/miguno/a9b44c1efc9b579565cfad75e543c808>

Configurations , start and stop

```
final Serde<String> stringSerde = Serdes.String();
final Serde<Long> longSerde = Serdes.Long();

KStreamBuilder builder = ...;
Properties streamsConfiguration = new Properties();
streamsConfiguration.put(StreamsConfig.APPLICATION_ID_CONFIG, "wordcount-lambda-example");
streamsConfiguration.put(StreamsConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
streamsConfiguration.put(StreamsConfig.ZOOKEEPER_CONNECT_CONFIG, "localhost:2181");
streamsConfiguration.put(StreamsConfig.KEY_SERDE_CLASS_CONFIG, Serdes.String().getClass().getName());
streamsConfiguration.put(StreamsConfig.VALUE_SERDE_CLASS_CONFIG, Serdes.String().getClass().getName());
// Requires Kafka 0.10.1 or later.

// Upcoming Kafka 0.10.1 does this automatically for you (no need for through).
.flatMapValues((key, value) -> {
    .map((key, value) -> {
        // ...
        .through("RekeyedIntermediateTopic")
        .countByKey("Counts")
        .toStream();
    });
    wordCounts.to(stringSerde, longSerde, "WordsWithCountsTopic");
});

KafkaStreams streams = new KafkaStreams(builder, streamsConfiguration);
streams.start();

Runtime.getRuntime().addShutdownHook(new Thread(streams::close));
```



<https://gist.github.com/miguno/a9b44c1efc9b579565cfad75e543c808>

```

public class WordCountProcessor implements Processor<String, String> {

    private ProcessorContext context;
    private KeyValueStore<String, Long> kvStore;

    @Override
    @SuppressWarnings("unchecked")
    public void init(ProcessorContext context) {
        // keep the processor context locally because we need it in punctuate() and commit()
        this.context = context;

        // call this processor's punctuate() method every 1000 time units.
        this.context.schedule(1000);

        // retrieve the key-value store named "Counts"
        kvStore = (KeyValueStore) context.getStateStore("Counts");
    }

    @Override
    public void process(String dummy, String line) {
        String[] words = line.toLowerCase().split(" ");

        for (String word : words) {
            Long oldValue = kvStore.get(word);
            if (oldValue == null) {
                kvStore.put(word, 1L);
            } else {
                kvStore.put(word, oldValue + 1L);
            }
        }
    }

    @Override
    public void punctuate(long timestamp) {
        KeyValueIterator<String, Long> iter = this.kvStore.all();
        while (iter.hasNext()) {
            KeyValue<String, Long> entry = iter.next();
            context.forward(entry.key, entry.value.toString());
        }
        iter.close();
        // commit the current processing progress
        context.commit();
    }

    @Override
    public void close() {
        // close the key-value store
        kvStore.close();
    }
}

```

Kafka Streams for Stream processing

<https://balamaci.ro/kafka-streams-for-stream-processing/>

<https://docs.confluent.io/3.2.0streams/developer-guide.html#processor-api>

processor

@Override

init(ProcessorContext context)

@Override

process(String key, String value)

@Override

punctuate(long timestamp)

@Override

close()

punctuate method will be executed periodically - with the frequency we set in **init()** call to `context.schedule(1000)`. We send the data to the next Processor with **context.forward(entry.key, entry.value.toString())**.



State stores

RockDB

```
// Creating a persistent key-value store:  
// here, we create a `KeyValueStore<String, Long>`  
import org.apache.kafka.streams.processor.StateStore  
import org.apache.kafka.streams.state.Stores;  
  
// Note: The `Stores` factory returns a supplier for  
// because that's what you typically need to pass as  
StateStoreSupplier countStoreSupplier =  
    Stores.create("persistent-counts")  
        .withKeys(Serdes.String())  
        .withValues(Serdes.Long())  
        .persistent()  
        .build();
```

In memory

```
// Creating an in-memory key-value store:  
// here, we create a `KeyValueStore<String, Long>`  
import org.apache.kafka.streams.processor.StateStore  
import org.apache.kafka.streams.state.Stores;  
  
// Note: The `Stores` factory returns a supplier for  
// because that's what you typically need to pass as  
StateStoreSupplier countStoreSupplier =  
    Stores.create("inmemory-counts")  
        .withKeys(Serdes.String())  
        .withValues(Serdes.Long())  
        .inMemory()  
        .build();
```

```
import org.apache.kafka.streams.processor.StateStoreSupplier;  
import org.apache.kafka.streams.state.Stores;  
  
Map<String, String> changelogConfig = new HashMap();  
// override min.insync.replicas  
changelogConfig.put("min.insync.replicas", "1")  
  
StateStoreSupplier countStoreSupplier = Stores.create("Counts")  
    .withKeys(Serdes.String())  
    .withValues(Serdes.Long())  
    .persistent()  
    .enableLogging(changelogConfig) // enable changelogging, with custom changelog settings  
    .build();
```

Connecting Processors and State Stores

```
TopologyBuilder builder = new TopologyBuilder();

// add the source processor node that takes Kafka topic "source-topic" as input
builder.addSource("Source", "source-topic")

// add the WordCountProcessor node which takes the source processor as its upstream processor
.addProcessor("Process", () -> new WordCountProcessor(), "Source")

// add the count store associated with the WordCountProcessor processor
.addStateStore(countStoreSupplier, "Process")

// add the sink processor node that takes Kafka topic "sink-topic" as output
// and the WordCountProcessor node as its upstream processor
.addSink("Sink", "sink-topic", "Process");
```

Kafka Streams for Stream processing
<https://balamaci.ro/kafka-streams-for-stream-processing/>

State store

- DSL API restricts access to a read-only view of the state store (via [ReadOnlyKeyValueStore](#)) as opposed to the Processor API using which you can get access to a writable view which allows you to mutate the state store

```
//DSL API
```

```
KafkaStreams ks = ...;
ReadOnlyKeyValueStore<String, Double> myStore = ks.store(storeName,
                                                       QueryableStoreTypes.<String, Double>keyValueStore());
```

```
//Processor API
```

```
ProcessorContext pc = ... ;
KeyValueStore<String, Double> kvStateStore = (KeyValueStore<String, Double>) pc.getStateStore("my-state-store");
kvStateStore.put("my-key",2); //mutate operation
```

State store

```
public class MyProductCountProcessor extends AbstractProcessor<Long, String> {  
    private ProcessorContext context;  
    private KeyValueStore<Long, Long> kvStore;  
  
    public void init(ProcessorContext context) {  
        (...)  
        this.kvStore = (KeyValueStore<Long, Long>) context.getStateStore("Counts");  
    }  
  
    public void process(Long productId, String browser) {  
        Long oldValue = this.kvStore.get(word);  
        if (oldValue == null) {  
            this.kvStore.put(word, 1L);  
        } else {  
            this.kvStore.put(word, oldValue + 1L);  
        }  
    }  
  
    //This method will be called as a way to do something periodically  
    public void punctuate(long timestamp) {  
        KeyValueIterator<Long, Long> iter = this.kvStore.all();  
  
        while (iter.hasNext()) {  
            KeyValue<String, Long> entry = iter.next();  
            context.forward(entry.key, entry.value.toString());  
        }  
        iter.close();  
        context.commit();  
    }  
}
```

Kafka Streams for Stream processing
<https://balamaci.ro/kafka-streams-for-stream-processing/>

GlobalKTable

A GlobalKTable is created via a KStreamBuilder. For example:

```
builder.globalTable("topic-name", "queryable-store-name");
```

all GlobalKTables are backed by a ReadOnlyKeyValueStore and are therefore queryable via the interactive queries API. For example:

```
final GlobalKTable globalOne = builder.globalTable("g1", "g1-store");
final GlobalKTable globalTwo = builder.globalTable("g2", "g2-store");
...
final KafkaStreams streams = ...;
streams.start()
...
ReadOnlyKeyValueStore view = streams.store("g1-store",
QueryableStoreTypes.keyValueStore());
view.get(key); // can be done on any key, as all keys are present
```

Kafka Streams for Stream processing

<https://balamaci.ro/kafka-streams-for-stream-processing/>

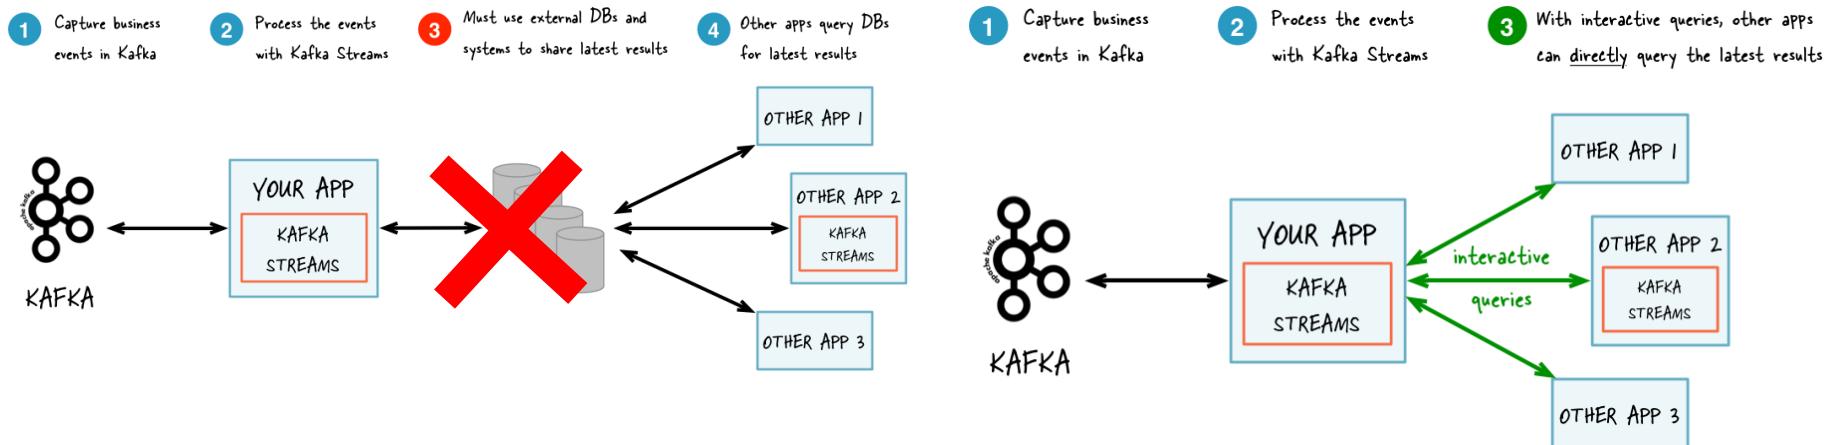
<https://github.com/confluentinc/kafka-streams-examples/blob/4.0.0-post/src/main/java/io/confluent/examples/streams/GlobalKTablesExample.java>

<https://kafka.apache.org/0102/javadoc/org/apache/kafka/streams/kstream/GlobalKTable.html>

Interactive Queries

no

with



Interactive Queries

```
StreamsConfig config = ...;
KStreamBuilder builder = ...;
KStream<String, String> textLines = ...;

// Define the processing topology (here: WordCount)
KGroupedStream<String, String> groupedByWord = textLines
    .flatMapValues(value -> Arrays.asList(value.toLowerCase().split("\\\\W+")))
    .groupBy((key, word) -> word, stringSerde, stringSerde);

// Create a window state store named "CountsWindowStore" that contains the word counts for every minute
groupedByWord.count(TimeWindows.of(60000), "CountsWindowStore");
```

```
// Get the window store named "CountsWindowStore"
ReadOnlyWindowStore<String, Long> windowStore =
    streams.store("CountsWindowStore", QueryableStoreTypes.windowStore());

// Fetch values for the key "world" for all of the windows available in this application instance
// To get *all* available windows we fetch windows from the beginning of time until now.
long timeFrom = 0; // beginning of time = oldest available
long timeTo = System.currentTimeMillis(); // now (in processing-time)
WindowStoreIterator<Long> iterator = windowStore.fetch("world", timeFrom, timeTo);
while (iterator.hasNext()) {
    KeyValue<Long, Long> next = iterator.next();
    long windowTimestamp = next.key;
    System.out.println("Count of 'world' @ time " + windowTimestamp + " is " + next.value);
}
```

<https://kafka.apache.org/10/documentationstreams/developer-guide/interactive-queries.html>

```

static WordCountInteractiveQueriesRestService startRestProxy(final KafkaStreams streams, final int port)
    throws Exception {
    final HostInfo hostInfo = new HostInfo(DEFAULT_HOST, port);
    final WordCountInteractiveQueriesRestService
        wordCountInteractiveQueriesRestService = new WordCountInteractiveQueriesRestService(streams, hostInfo);
    wordCountInteractiveQueriesRestService.start(port);
    return wordCountInteractiveQueriesRestService;
}

```

REST interface

```

static KafkaStreams createStreams(final Properties streamsConfiguration) {
    final Serde<String> stringSerde = Serdes.String();
    StreamsBuilder builder = new StreamsBuilder();
    KStream<String, String>
        textLines = builder.stream(TEXT_LINES_TOPIC, Consumed.with(Serdes.String(), Serdes.String()));

    final KGroupedStream<String, String> groupedByWord = textLines
        .flatMapValues(value -> Arrays.asList(value.toLowerCase().split("\\W+")))
        .groupBy((key, word) -> word, Serialized.with(stringSerde, stringSerde));

    // Create a State Store for with the all time word count
    groupedByWord.count(Materialized.<String, Long, KeyValueStore<Bytes, byte[]>>as("word-count")
        .withValueSerde(Serdes.Long()));

    // Create a Windowed State Store that contains the word count for every
    // 1 minute
    groupedByWord.windowedBy(TimeWindows.of(60000))
        .count(Materialized.<String, Long, WindowStore<Bytes, byte[]>>as("windowed-word-count")
        .withValueSerde(Serdes.Long()));

    return new KafkaStreams(builder.build(), streamsConfiguration);
}

```

<https://github.com/confluentinc/kafka-streams-examples/tree/4.0.0-post/src/main/java/io/confluent/examplesstreams/interactivequeries>

```

static WordCountInteractiveQueriesRestService startRestProxy(final KafkaStreams streams, final int port)
    throws Exception {
    final HostInfo hostInfo = new HostInfo(DEFAULT_HOST, port);
    final WordCountInteractiveQueriesRestService
        wordCountInteractiveQueriesRestService = new WordCountInteractiveQueriesRestService(streams, hostInfo);
    wordCountInteractiveQueriesRestService
    return wordCountInteractiveQueriesRestService;
}

static KafkaStreams createStreams() {
    final Serde<String> stringSerde = Serdes.String();
    StreamsBuilder builder = new StreamsBuilder();
    KStream<String, String>
        textLines = builder.stream("text-lines");
        .mapValues((value, key) -> value.toLowerCase());
        .groupByKey();
    final KGroupedStream<String, String>
        .flatMapValues((value, key) -> {
            final Map<String, Integer> wordCounts = new HashMap<>();
            for (final String word : value) {
                wordCounts.put(word, wordCounts.getOrDefault(word, 0) + 1);
            }
            return wordCounts.entrySet();
        })
        .groupByKey();
    // Create a State Store for word counts
    groupedByWord.count(Materialized.Value)
        .withValueSerde(Serdes.Integer());
    // Create a Windowed State for windowed word counts
    // 1 minute
    groupedByWord.windowedBy(TimerWindows.of(Duration.ofMinutes(1)))
        .count(Materialized.Value)
        .withValueSerde(Serdes.Long());
}

return new KafkaStreams(builder.build(), streamsConfiguration);
}

```

REST interface

The example is using JAX-RS over jetty

Get store view (queriable interface)

```

// Lookup the WindowStore with the provided storeName
final ReadOnlyWindowStore<String, Long> store = streams.store(storeName,
    QueryableStoreTypes.<String, Long>windowStore());
if (store == null) {
    throw new NotFoundException();
}

// fetch the window results for the given key and time range
final WindowStoreIterator<Long> results = store.fetch(key, from, to);

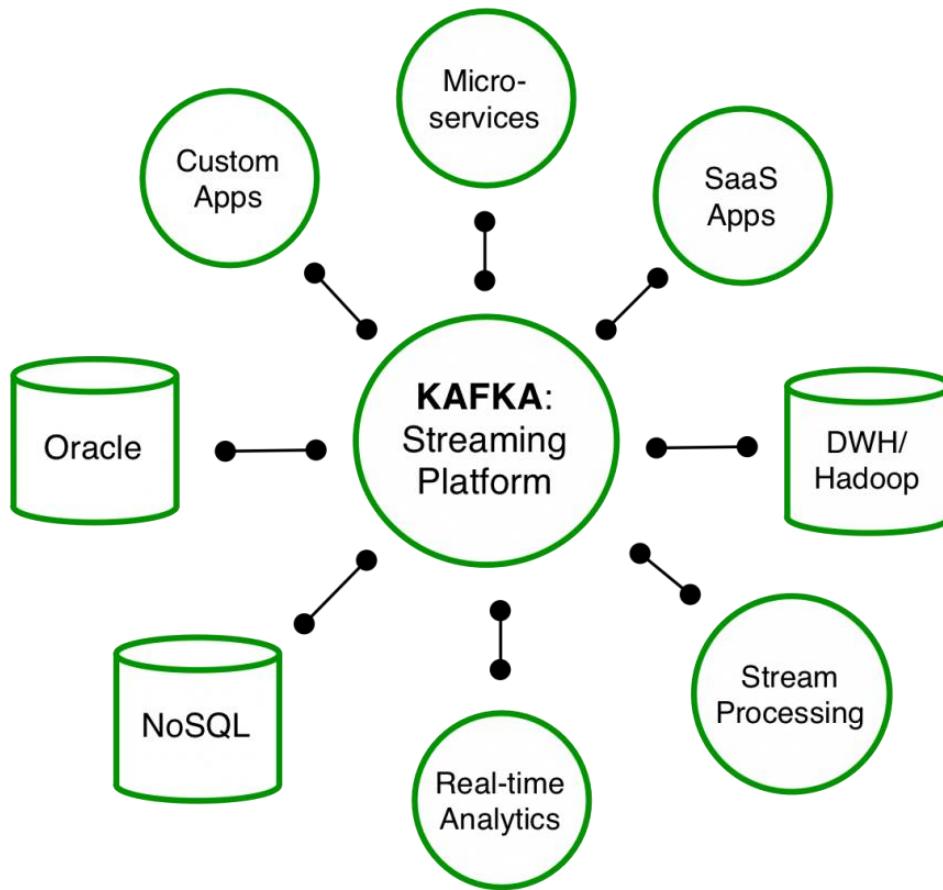
final List<KeyValueBean> windowResults = new ArrayList<>();
while (results.hasNext()) {
    final KeyValue<Long, Long> next = results.next();
    // convert the result to have the window time and the key (for display purposes)
    windowResults.add(new KeyValueBean(key + "@" + next.key, next.value));
}
return windowResults;
}

```

Expose data for REST

<https://github.com/confluentinc/kafka-streams-examples/tree/4.0.0-post/src/main/java/io/confluent/examplesstreams/interactivequeries>



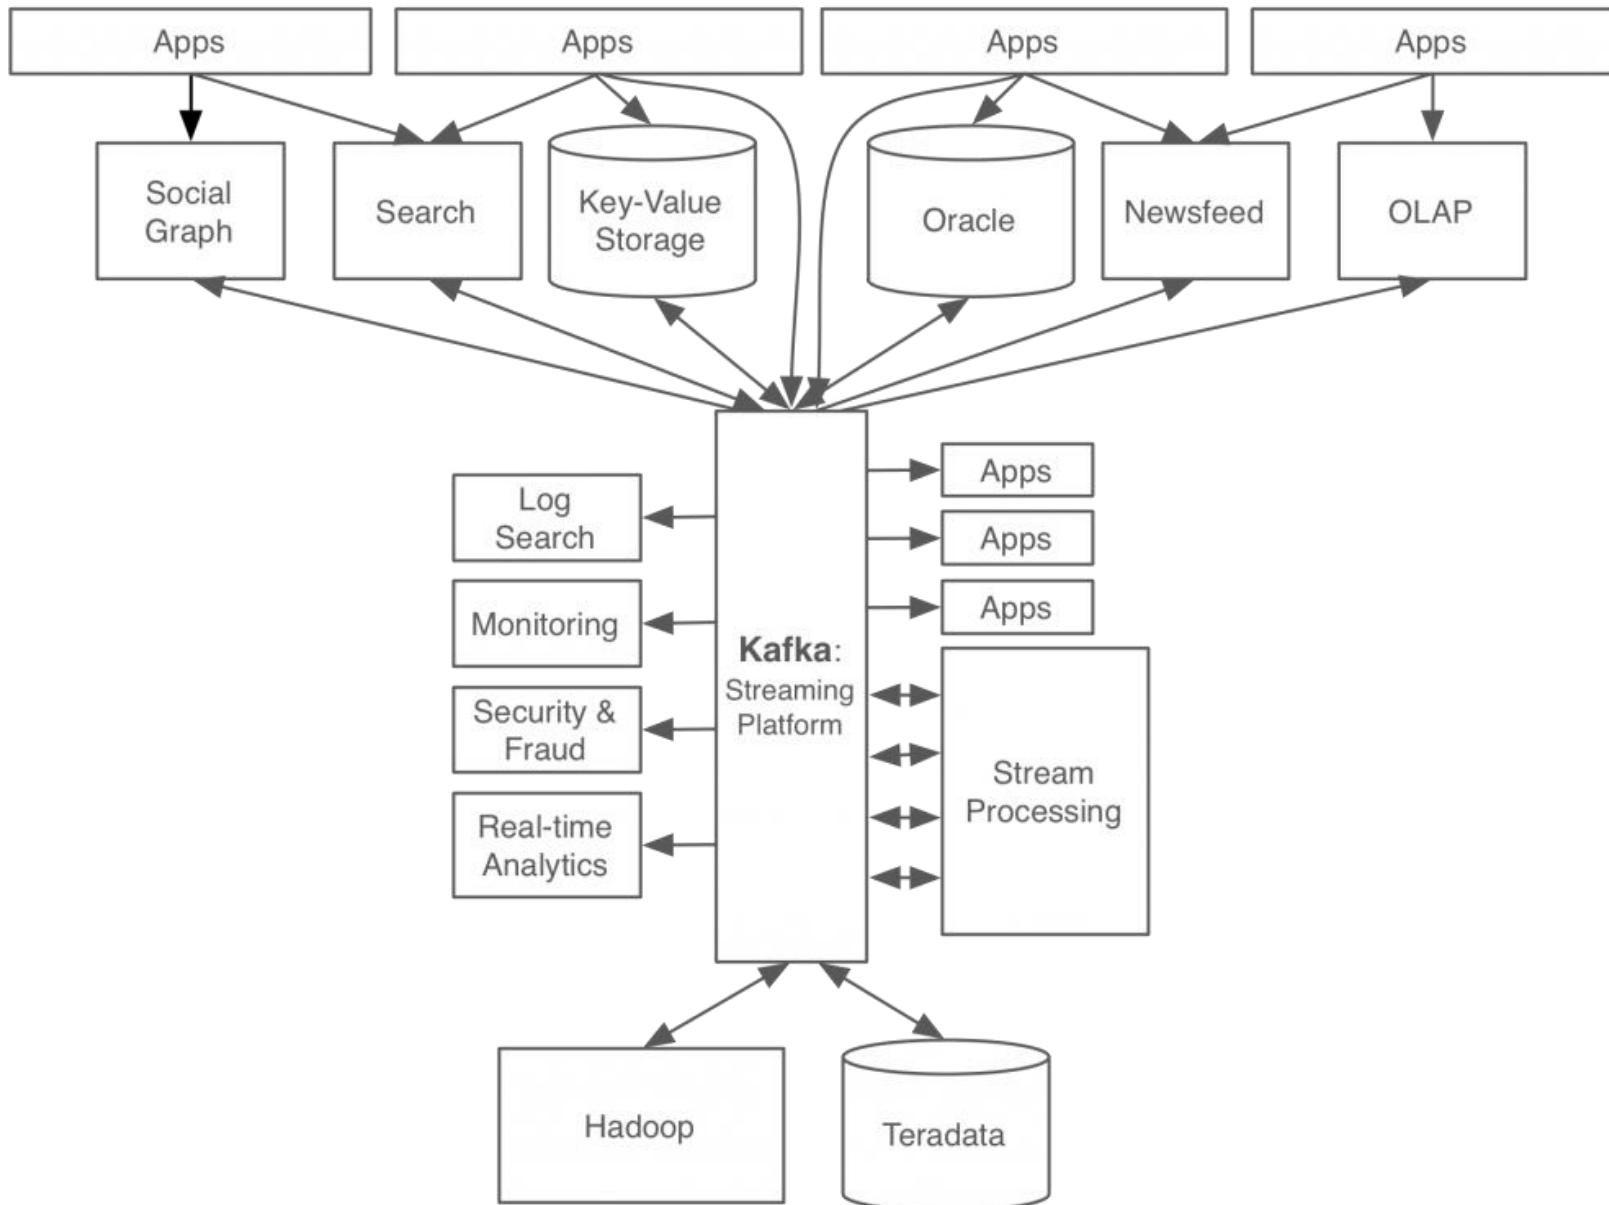


**Synchronous
Req/Response**
0 – 100s ms

Near real time
 $> 100s$ ms

Offline batch
 > 1 hour

<https://www.confluent.io/blog/stream-data-platform-1/>



<https://www.confluent.io/blog/stream-data-platform-1/>

Spring cloud

- “Spring Cloud Stream is a framework built on top of Spring Boot and Spring Integration that helps in creating event-driven or message-driven microservices.”

```
@SpringBootApplication
@EnableBinding(Processor.class)
public class MyLoggerServiceApplication {
    public static void main(String[] args) {
        SpringApplication.run(MyLoggerServiceApplication.class, args);
    }

    @StreamListener(Processor.INPUT)
    @SendTo(Processor.OUTPUT)
    public LogMessage enrichLogMessage(LogMessage log) {
        return new LogMessage(String.format("[1]: %s", log.getMessage()));
    }
}
```

Introduction to Spring Cloud Stream
<http://www.baeldung.com/spring-cloud-stream>

Spring cloud kafka streaming binding

```
@SpringBootApplication
@EnableBinding(KStreamProcessor.class)
public class WordCountProcessorApplication {

    @StreamListener("input")
    @SendTo("output")
    public KStream<?, String> process(KStream<?, String> input) {
        return input
            .flatMapValues(value -> Arrays.asList(value.toLowerCase().split("\\W+")))
            .map((key, word) -> new KeyValue<>(word, word))
            .groupByKey(Serdes.String(), Serdes.String())
            .count(TimeWindows.of(5000), "store-name")
            .toStream()
            .map((w, c) -> new KeyValue<>(null, "Count for " + w.key() + ": " + c));
    }

    public static void main(String[] args) {
        SpringApplication.run(WordCountProcessorApplication.class, args);
    }
}
```

Spring Cloud Stream Reference Guide

<https://docs.spring.io/spring-cloud-stream/docs/current/reference/htmlsingle/>

PROJECTS : SPRING CLOUD

Spring Cloud Data Flow



Spring Cloud Data Flow is a toolkit for building data integration and real-time data processing pipelines.

Pipelines consist of Spring Boot apps, built using the Spring Cloud Stream or Spring Cloud Task microservice frameworks. This makes Spring Cloud Data Flow suitable for a range of data processing use cases, from import/export to event streaming and predictive analytics.

QUICK START

<https://cloud.spring.io/spring-cloud-dataflow/#quick-start>

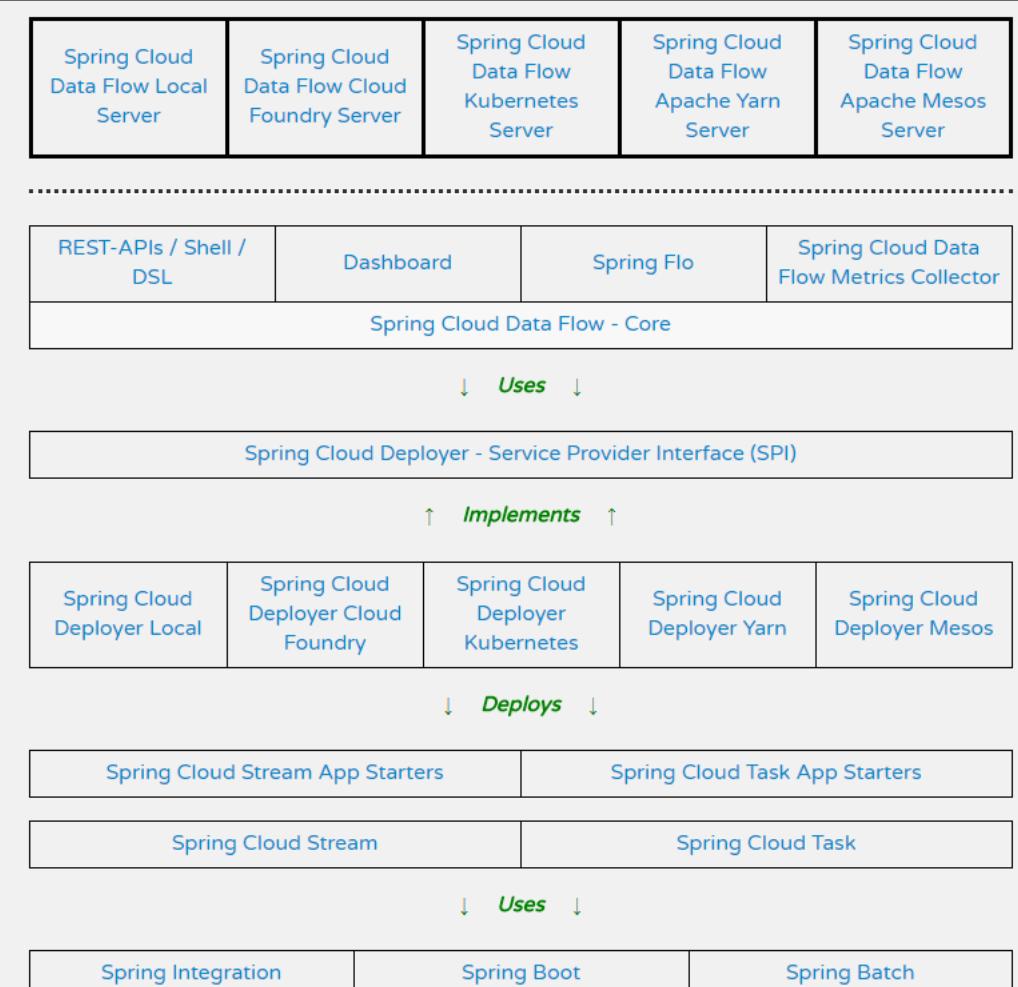
PROJECTS : SPRING CLOUD

Spring Cloud Data Flow

Spring Cloud Data Flow is a tool for building distributed data processing pipelines.

Pipelines consist of Spring Boot Task microservice frameworks for various data processing use cases, from

QUICK START



PROJECTS : SPRING CLOUD

Spring Cloud Data Flow

Spring Cloud Data Flow is a tool for building distributed data processing pipelines.

Pipelines consist of Spring Boot Task microservice frameworks for various data processing use cases, from

QUICK START

Spring Cloud Data Flow Local Server	Spring Cloud Data Flow Cloud Foundry Server	Spring Cloud Data Flow Kubernetes Server	Spring Cloud Data Flow Apache Yarn Server	Spring Cloud Data Flow Apache Mesos Server
-------------------------------------	---	--	---	--

REST-APIs / Shell / DSL	Dashboard	Spring Flo	Spring Cloud Data Flow Metrics Collector
Spring Cloud Data Flow - Core			

↓ *Uses* ↓

Spring Cloud Deployer - Service Provider Interface (SPI)
--

↑ *Implements* ↑

Spring Cloud Deployer Local	Spring Cloud Deployer Cloud Foundry	Spring Cloud Deployer Kubernetes	Spring Cloud Deployer Yarn	Spring Cloud Deployer Mesos
-----------------------------	-------------------------------------	----------------------------------	----------------------------	-----------------------------

↓ *Deploys* ↓

Spring Cloud Stream App Starters	Spring Cloud Task App Starters
Spring Cloud Stream	Spring Cloud Task

↓ *Uses* ↓

Spring Integration	Spring Boot	Spring Batch
--------------------	-------------	--------------

<https://cloud.spring.io/spring-cloud-dataflow/#quick-start>

<https://cloud.spring.io/spring-cloud-stream-app-starters/>



PROJECTS : SPRING CLOUD

Spring Cloud Data Flow

Spring Cloud Data Flow is a tool for building distributed data processing pipelines.

Pipelines consist of Spring Boot Task microservice frameworks for various data processing use cases, from

QUICK START

<https://cloud.spring.io/spring-cloud-dataflow/>
<https://cloud.spring.io/spring-cloud-stream/>

41492- Engenharia de Software / Software Engineering

Source	Processor	Sink
file	aggregator	aggregate-counter
ftp	bridge	cassandra
gemfire	filter	counter
gemfire-cq	groovy-filter	field-value-counter
http	groovy-transform	file
jdbc	header-enricher	ftp
jms	httpclient	gemfire
load-generator	pmmi	gpfldist
loggregator	python-http	hdfs
mail	python-jython	hdfs-dataset
mongodb	scriptable-transform	jdbc
mqtt	splitter	log
rabbit	tasklauchrequest-transform	mongodb
s3	tcp-client	mqtt
sftp	tensorflow	pgcopy
syslog	transform	rabbit
tcp	twitter-sentiment	redis-pubsub
tcp-client		router
time		s3
trigger		sftp
triggertask		task-launcher-cloudfoundry
twitterstream		task-launcher-local
		task-launcher-yarn
		tcp
		throughput
		websocket



KSQL

Streaming SQL for Apache Kafka

KSQL is an open source, Apache 2.0 licensed streaming SQL engine that enables stream processing against Apache Kafka®.

KSQL makes it easy to read, write, and process streaming data in real-time, at scale, using SQL-like semantics. It offers an easy way to express stream processing transformations as an alternative to writing an application in a programming language such as Java or Python.

KSQL provides powerful stream processing capabilities such as joins, aggregations, event-time windowing, and more!



Try KSQL

[Notify me when KSQL is generally available](#)



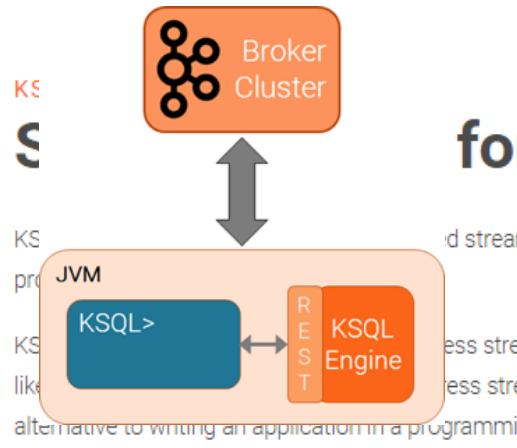
KSQL: Query your streams without writing code

<https://www.confluent.io/product/ksql/>

<https://github.com/confluentinc/ksql/blob/v0.4/docs/concepts.md#concepts>

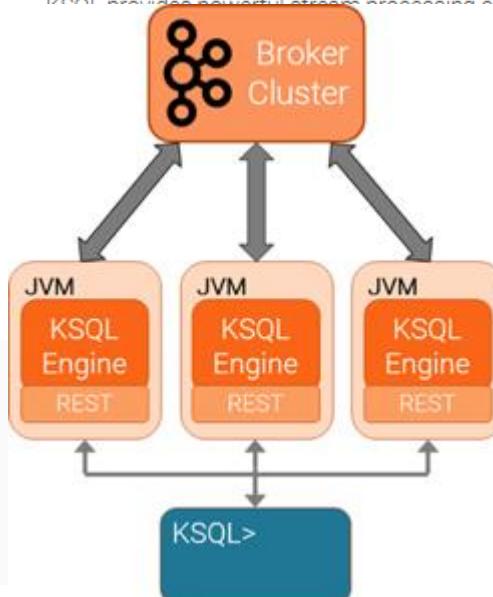
ssing against Kafka today.





```
CREATE STREAM vip_actions AS  
SELECT userid, page, action  
FROM clickstream c  
LEFT JOIN users u ON c.userid = u.user_id  
WHERE u.level = 'Platinum';
```

KSQL provides powerful stream processing capabilities such as joins, aggregations, event-time



```
CREATE TABLE error_counts AS  
SELECT error_code, count(*)  
FROM monitoring_stream  
WINDOW TUMBLING (SIZE 1 MINUTE)  
WHERE type = 'ERROR'  
GROUP BY error_code;
```

try your streams without writing code

missing against Kafka today.



LANDOOP

PRODUCTS



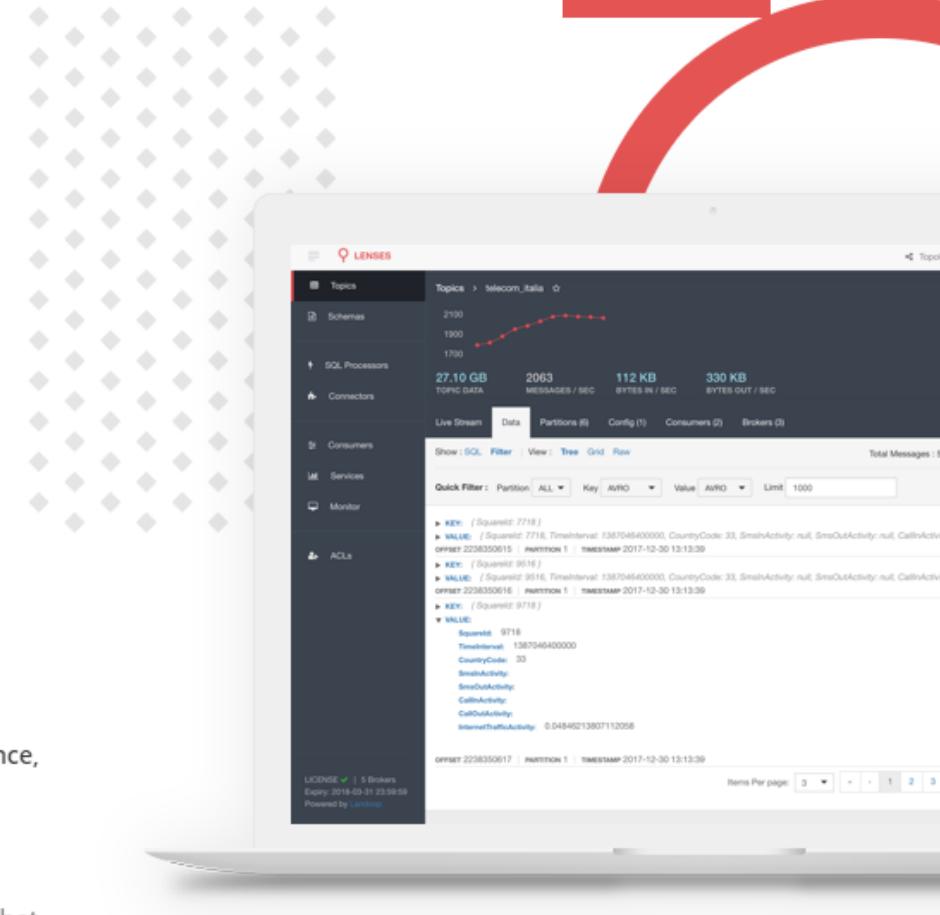
Successful streaming over Apache Kafka®

- ✓ Demystify complexity, Focus on what's important.
- ✓ Optimise engineering resources
- ✓ Robust, multi-level production tooling for extreme resilience, at scale

Lenses is an enterprise grade product that provides faster streaming application deliveries and data flow management that natively integrates over Apache Kafka.

<http://www.landoop.com/kafka-lenses/>

<http://www.landoop.com/kafka/kafka-sql/> Kafka with a rich user



Got it!



universidade de aveiro

LANDOOP
PRODU



Successful st

Kafka ®

- Demystify complex data
- Optimise engineering
- Robust, multi-level at scale

Lenses is an enterprise streaming application that natively integrates with

<http://www.landoop.com>
<http://www.landoop.com>



CONNECT



ANALYSE



PROCESS

Analyse

Data Browsing, Topics and Messages, Choose Decoders, AVRO supported

Lenses SQL knows how to handle the following payload structures (also known as decoder types): BINARY, JSON, AVRO, STRING, INT, LONG, PROTOBUF*(coming soon).

AVRO Schema Registry is also fully supported.

```
SELECT field2 AS my_field
FROM magicTopic
WHERE _ktype=DOUBLE
    AND _vtype=AVRO
```

```
SELECT *
FROM magicTopic
WHERE _ktype=JSON
    AND _vtype=AVRO
```

Access Metadata

```
SELECT
    field1,
    _offset,
    _ts,
    _partition
FROM magicTopic
```

Work with Keys

```
SELECT min_temp AS 'min'
FROM sensorsTopic
WHERE _key LIKE 'RN81'
    AND _ktype=STRING
    AND _vtype=AVRO
    AND _partition = 0
```

Nested Fields

```
SELECT
    details.antennaclass AS antenna,
    details.readingperiod AS 'interval'
FROM sensorsTopic
```

Filtering

```
SELECT *
FROM topicA
WHERE (a.d.e + b) / c = 100
    AND _offset > 2000
    AND partition IN (2,6,9)
LIMIT 1000
```

Project Fields & Functions

```
SELECT field2 * field3 - abs(field4.field5) AS total,
    _offset AS '_key_offset'
FROM magicTopic
```

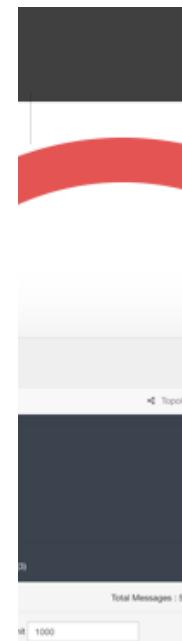
Time Travelling

```
SELECT
    device_id,
    temperature
FROM sensorsTopic
WHERE _ts > '2017-08-01 12:24:00' AND _ts < '2017-08-01 14:30:00'
```

```
SELECT min_temp AS 'min',
    max_temp AS 'max',
    (min_temp + max_temp) / 2 AS average,
    details.antennaclass AS antenna,
    details.readingperiod AS 'interval'
FROM sensorsTopic
WHERE _key_id like 'RN81'
    AND _ktype=JSON
    AND _vtype=JSON
```

< CONNECT

PROCESS >



Our website uses some cookies to provide you with a better experience, user-functionality, extra features and help us with troubleshooting.

Got it!



LANDO

```
SET `auto.create`=true;
SET `auto.offset.reset`='earliest';
SET `commit.interval.ms`='120000';
SET `compression.type`='snappy';
```

```
INSERT INTO `cc_payments_fraud`
WITH tableCards AS (
    SELECT *
    FROM `cc_data`
    WHERE _ktype='STRING' AND _vtype='AVRO' )
```

```
SELECT STREAM
    p.currency,
    sum(p.amount) as total,
    count(*) usage
FROM `cc_payments` AS p LEFT JOIN tableCards AS c ON p._key = c._key
WHERE p._ktype='STRING' AND p._vtype='AVRO' and c.blocked is true
GROUP BY tumble(1,m), p.currency
```

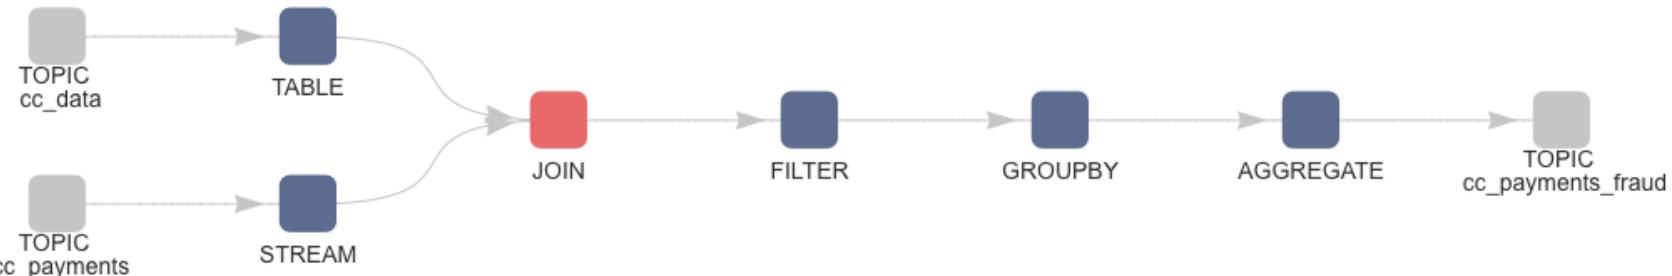
Successful

Kafka®

Demystify complexity, Focus on what's important.

Stream Topology

click on the nodes for details

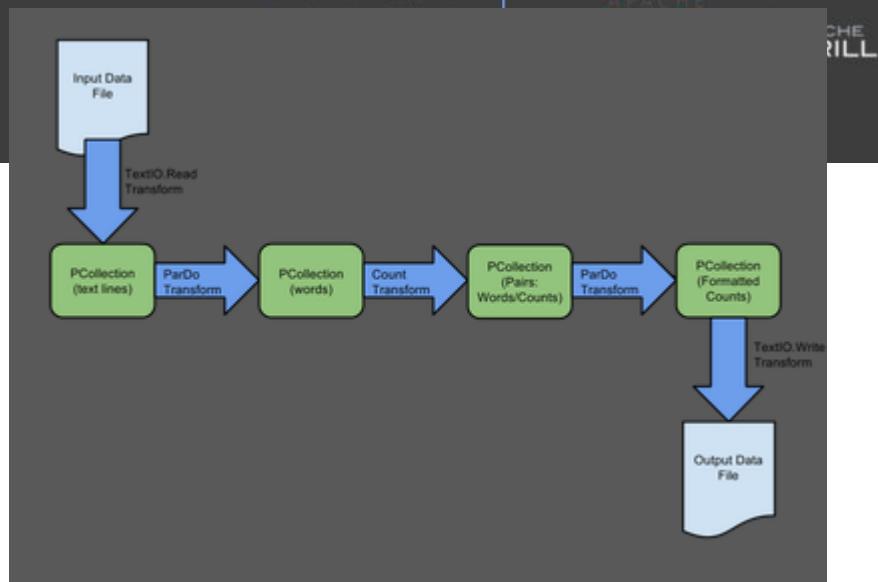


Our website uses some cookies to provide you with a better experience, user-functionality, extra features and help us with troubleshooting.

Got it!



The Evolution of Apache Beam



```
Pipeline p = Pipeline.create(options);
p.apply(TextIO.Read.from("...kinglear.txt/input for this transform
is the PCollection of text lines
//generated by the TextIO.Read transform
// Anonymous DoFn executed on each element that tokenizes the text
lines into individual words.
.apply(ParDo.named("ExtractWords").of(newng, String>() {
    @Override
    public void processElement(ProcessContext c) {
        for (String word : c.element().split("[^a-zA-Z]+")) {
            if (!word.isEmpty()) {
                c.output(word);
            }
        }
    }
))
// each element in output PCollection represents an individual
word in the text.
```

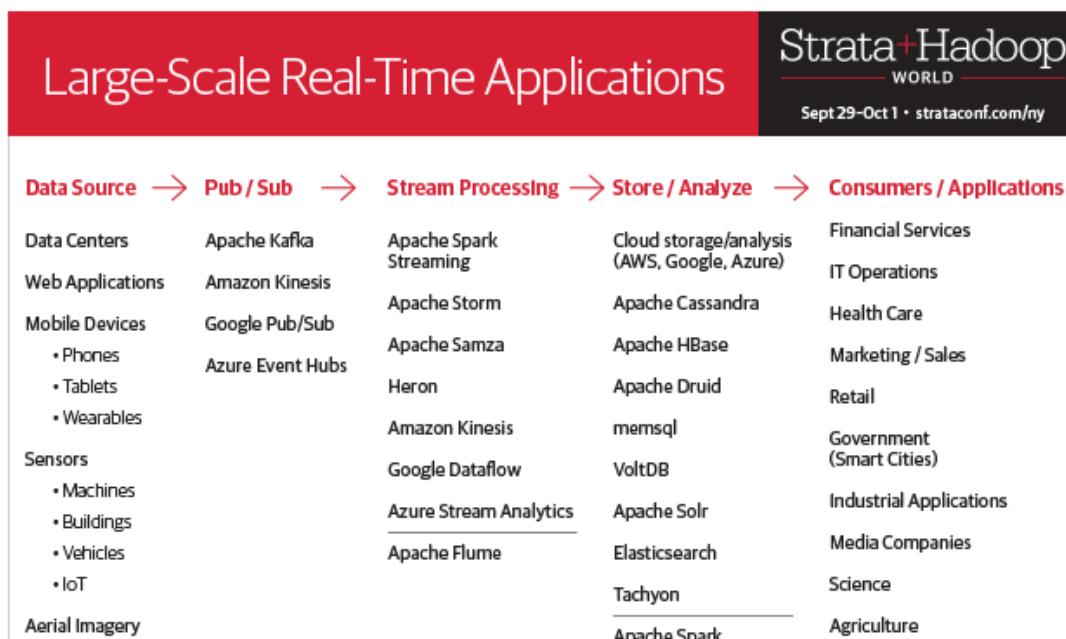
Beam

What?	Sums of integers, keyed by team.
Where?	Within fixed event-time windows of one hour.
When?	<ul style="list-style-type: none">• Early: Every 5 minutes of processing time.• On-time: When the watermark passes the end of the window.• Late: Every 10 minutes of processing time.• Final: When the watermark passes the end of the window + two hours.
How?	Panes accumulate new values into prior results.

```
gameEvents
    [... input ...]
    .apply("LeaderboardTeamFixedWindows", Window
        .<GameActionInfo>into(FixedWindows.of(
            Duration.standardMinutes(Durations.minutes(60))))
        .triggering(AfterWatermark.pastEndOfWindow()
            .withEarlyFirings(AfterProcessingTime.pastFirstElementInPane()
                .plusDelayOf(Durations.minutes(5)))
            .withLateFirings(AfterProcessingTime.pastFirstElementInPane()
                .plusDelayOf(Durations.minutes(10))))
            .withAllowedLateness(Duration.standardMinutes(120))
            .accumulatingFiredPanes()))
        .apply("ExtractTeamScore", new ExtractAndSumScore("team"))
    [... output ...]
```

ANALYTICS

Real-time Streaming ETL and Real-time BI – Part 2



Real-time Streaming ETL and Real-time BI – Part 2

<https://phismmetry.wordpress.com/2015/09/26/real-time-streaming-etl-and-real-time-bi-part-2/>

TECHNOLOGY MUSINGS

Just another WordPress.com site

ANALYTICS

Real-time Analytics

Large
Data

Data Source

Data Centers

Web Application

Mobile Devices

• Phones

• Tablets

• Wearables

Sensors

• Machines

• Buildings

• Vehicles

• IoT

Aerial Imagery

Now , how does the overall stack look like !

- [Zookeeper](#) handles much of the system coordination (via Service Managers like [Helix](#) or [Curator](#)).
- [Mesos](#) and [YARN](#) process virtualization and resource management.
- Embedded libraries like [Lucene](#), [RocksDB](#), [FastBits](#) and [LMDB](#) do indexing.
- [Netty](#), [Jetty](#), and higher-level wrappers like [Finagle](#) and [rest.li](#) handle remote communication.
- [Avro](#), [Protocol Buffers](#), [Thrift](#), and [umpteen zillion](#) other libraries handle serialization.
- [Kafka](#) and [BookKeeper](#) provide a backing log.

Azure Event Hubs	Heron	Apache Druid	Retail
	Amazon Kinesis	memsql	Government (Smart Cities)
	Google Dataflow	VoltDB	Industrial Applications
	Azure Stream Analytics	Apache Solr	Media Companies
	Apache Flume	Elasticsearch	Science
		Tachyon	Agriculture
		Apache Spark	

Real-time Streaming ETL and Real-time BI – Part 2

<https://phismmetry.wordpress.com/2015/09/26/real-time-streaming-etl-and-real-time-bi-part-2/>

The END